

Automobile Price Prediction

About Dataset This dataset consist of data From 1985 Ward's Automotive Yearbook.

Content This data set consists of three types of entities:

- (a) the specification of an auto in terms of various characteristics,
- (b) its assigned insurance risk rating,
- (c) its normalized losses in use as compared to other cars. The second rating corresponds to the degree to which the auto is more risky than its price indicates. Cars are initially assigned a risk factor symbol associated with its price. Then, if it is more risky (or less), this symbol is adjusted by moving it up (or down) the scale. Actuarians call this process "symboling". A value of +3 indicates that the auto is risky, -3 that it is probably pretty safe.
- The third factor is the relative average loss payment per insured vehicle year. This value is normalized for all autos within a particular size classification (two-door small, station wagons, sports/speciality, etc...), and represents the average loss per car per year.

Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msn

import warnings
warnings.filterwarnings('ignore')


pd.set_option('display.max_columns', None)
```

Load Dataset

```
In [2]: df = pd.read_csv('D:/DS Bootcamp/Machine Learning/machine learning projects/Machine-Learning-Projects/Automobile Price Prediction.csv')
df.sample(5)
```

Out[2]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width	height
14285	3.0	?	mitsubishi	gas	turbo	NaN	hatchback	fwd	front	95.9	173.2	66.3	50.2
23006	1.0	128	nissan	gas	std	two	hatchback	fwd	front	94.5	NaN	63.8	NaN
19921	0.0	?	jaguar	gas	std	four	sedan	rwd	front	113.0	199.6	69.6	52.8
1445	2.0	NaN	toyota	NaN	std	two	NaN	rwd	front	98.4	176.2	65.6	NaN
23744	0.0	161	peugot	gas	std	four	sedan	rwd	NaN	107.9	186.7	68.4	56.7



```
In [3]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30330 entries, 0 to 30329
Data columns (total 26 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   symboling              27286 non-null  float64
 1   normalized-losses      27294 non-null  object
 2   make                   27228 non-null  object
 3   fuel-type              27309 non-null  object
 4   aspiration              27355 non-null  object
 5   num-of-doors            27321 non-null  object
 6   body-style              27326 non-null  object
 7   drive-wheels            27215 non-null  object
 8   engine-location         27352 non-null  object
 9   wheel-base              27264 non-null  float64
10   length                 27258 non-null  float64
11   width                  27387 non-null  float64
12   height                  27281 non-null  float64
13   curb-weight             27302 non-null  float64
14   engine-type             27285 non-null  object
15   num-of-cylinders        27298 non-null  object
16   engine-size             27271 non-null  float64
17   fuel-system             27249 non-null  object
18   bore                   27373 non-null  object
19   stroke                  27409 non-null  object
20   compression-ratio       27327 non-null  float64
21   horsepower              27182 non-null  object
22   peak-rpm                27331 non-null  object
23   city-mpg                27229 non-null  float64
24   highway-mpg             27303 non-null  float64
25   price                   27287 non-null  object
dtypes: float64(10), object(16)
memory usage: 6.0+ MB

```

Handle Data Anomilies

```

In [4]: question_mark_counts = (df == '?').sum()
        question_mark_counts = question_mark_counts[question_mark_counts > 0]

```

```
print("Columns containing '?'\n")
print(question_mark_counts)
```

Columns containing '?'

```
normalized-losses    5526
num-of-doors         260
bore                 531
stroke              532
horsepower          260
peak-rpm            250
price               563
dtype: int64
```


Our first anomilie is '?' in the dataset. We will replace it with NaN

```
In [5]: for i in df[['normalized-losses', 'num-of-doors', 'bore', 'stroke', 'horsepower', 'peak-rpm', 'price']]:
        df[i].replace('?', np.nan, inplace = True)
```

Lets check '?' removed ore not

```
In [6]: df[df['bore'] == '?']
```

Out[6]:

symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width	height	curb-weight	en
														

FirsIt anomily solved

2nd we have wrong Dtypes of columns

- Change ['bore', 'stroke', 'normalized-losses', 'horsepower', 'peak-rpm', 'price'] into number

```
In [7]: for column in ['bore', 'stroke', 'normalized-losses', 'horsepower', 'peak-rpm', 'price']:
        df[column] = pd.to_numeric(df[column], errors='coerce')
```

```
In [8]: df.dtypes
```

```
Out[8]: symboling          float64
normalized-losses float64
make           object
fuel-type      object
aspiration     object
num-of-doors   object
body-style     object
drive-wheels   object
engine-location object
wheel-base    float64
length         float64
width          float64
height         float64
curb-weight    float64
engine-type    object
num-of-cylinders object
engine-size    float64
fuel-system    object
bore           float64
stroke        float64
compression-ratio float64
horsepower     float64
peak-rpm       float64
city-mpg       float64
highway-mpg    float64
price          float64
dtype: object
```

Data types issue resolved

Lets check more deeper in data

```
In [9]: print(f'Total no of numerical columns are: {len(df.select_dtypes(include='number').columns)}')
print('=====')
print(f'Total no of categorical columns are: {len(df.select_dtypes(include='object').columns)}')
print('=====')

for col in df.select_dtypes(include='number').columns:
    print(f'No of Unique values {col}: {df[col].nunique()}')
print('=====')
```

```
for col in df.select_dtypes(include='object').columns:
    print(f'\n{df[col].value_counts().head()}')
print('=====')

for col in df.select_dtypes(include='number').columns:
    print(f'\nUnique values of {col}: {df[col].unique()}')
```

```
Total no of numerical columns are: 16
=====
Total no of categorical columns are: 10
=====
No of Unique values symboling: 6
No of Unique values normalized-losses: 51
No of Unique values wheel-base: 53
No of Unique values length: 75
No of Unique values width: 44
No of Unique values height: 49
No of Unique values curb-weight: 171
No of Unique values engine-size: 44
No of Unique values bore: 38
No of Unique values stroke: 36
No of Unique values compression-ratio: 32
No of Unique values horsepower: 59
No of Unique values peak-rpm: 23
No of Unique values city-mpg: 29
No of Unique values highway-mpg: 30
No of Unique values price: 186
=====
```

```
make
toyota      4229
nissan       2338
mazda       2277
mitsubishi  1796
honda       1644
Name: count, dtype: int64
```

```
fuel-type
gas      24671
diesel   2638
Name: count, dtype: int64
```

```
aspiration
std      22383
turbo    4972
Name: count, dtype: int64
```

```
num-of-doors
four     15292
```

two 11769
Name: count, dtype: int64

body-style
sedan 12776
hatchback 9371
wagon 3361
hardtop 1016
convertible 802
Name: count, dtype: int64

drive-wheels
fwd 15777
rwd 10174
4wd 1264
Name: count, dtype: int64

engine-location
front 26990
rear 362
Name: count, dtype: int64

engine-type
ohc 19569
ohcf 2003
ohcv 1716
dohc 1671
l 1654
Name: count, dtype: int64

num-of-cylinders
four 21056
six 3310
five 1450
eight 672
two 521
Name: count, dtype: int64

fuel-system
mpfi 12506
2bbl 8834
idi 2622


```
1bb1      1405
spdi      1207
Name: count, dtype: int64
=====
```

```
Unique values of symboling: [ 3.  1.  2. nan  0. -1. -2.]
```

```
Unique values of normalized-losses: [ nan 164. 158. 192. 188. 121.  98.  81. 118. 148. 110. 145. 137. 101.
 78. 106.  85. 107. 104. 113. 150. 129. 115.  93. 161. 153. 125. 128.
103. 122. 108. 194. 231. 119. 154.  74. 186.  83. 102.  89.  87.  77.
 91. 168. 134.  65. 197.  90.  94. 256.  95. 142.]
```

```
Unique values of wheel-base: [ 88.6  94.5  99.8  99.4 105.8   nan  99.5 101.2 103.5 110.   88.4  93.7
103.3  95.9  86.6  96.5  94.3  96.   113.   102.   93.1  95.3  98.8 104.9
106.7 115.6  96.6 120.9 112.   102.7  93.   96.3  95.1  97.2 100.4  91.3
 99.2 107.9 114.2 108.   89.5  98.4  96.1  99.1  93.3  97.   96.9  95.7
102.4 102.9 104.5  97.3 104.3 109.1]
```

```
Unique values of length: [168.8 171.2 176.6   nan 177.3 192.7 178.2 176.8 189.   197.   141.1 155.9
157.3 174.6 144.6 150.   163.4 157.1 167.5 175.4 169.1 170.7 172.6 199.6
191.7 159.1 166.8 169.   177.8 175.   190.9 187.5 202.6 208.1 199.2 178.4
173.   173.2 172.4 165.3 170.2 165.6 162.4 173.4 181.7 184.6 186.7 198.9
167.3 168.9 175.7 181.5 186.6 157.9 172.   173.5 173.6 158.7 169.7 166.3
168.7 176.2 175.6 183.5 187.8 171.7 159.3 165.7 180.2 183.1 188.8 178.5
158.8 180.3 156.9 193.8]
```

```
Unique values of width: [64.1 65.5 66.2   nan 66.3 71.4 67.9 64.8 66.9 70.9 60.3 63.6 63.8 64.6
63.9 64.   65.2 66.   61.8 69.6 70.6 64.2 65.7 66.5 66.1 70.3 71.7 70.5
72.   68.   64.4 65.4 68.4 68.3 65.   72.3 66.6 63.4 65.6 67.7 67.2 68.8
68.9 66.4 62.5]
```

```
Unique values of height: [ nan 48.8 52.4 54.3 53.1 55.7 55.9 52.   53.7 56.3 53.2 50.8 50.6 59.8
50.2 52.6 54.5 58.3 53.3 54.1 51.   53.5 51.4 52.8 47.8 49.6 55.5 54.4
56.5 56.7 55.4 54.8 49.4 51.6 54.7 55.1 56.1 49.7 58.7 56.   50.5 55.2
52.5 53.   54.9 59.1 53.9 57.5 56.2 55.6]
```

```
Unique values of curb-weight: [2548. 2823. 2337. 2824. 2507. 2844. 2954. 3086. 3053. 2395. 2710. 2765.
3055. 3230. 3380. 3505.   nan 1874. 1909. 1876. 1989. 2191. 2811. 1713.
1819. 1837. 1940. 1956. 2010. 2024. 2236. 2289. 2304. 2372. 2465. 2293.
2734. 4066. 3950. 1890. 1900. 1905. 1945. 2380. 2385. 2500. 2410. 2443.
2425. 2670. 2700. 3515. 3750. 3495. 3770. 3740. 3685. 3900. 3715. 2910.
1944. 2004. 2370. 2328. 2833. 2921. 2405. 2403. 1889. 2017. 1938. 1951.]
```

2028. 1971. 2037. 2324. 2302. 3095. 3296. 3060. 3071. 3139. 3020. 3197.
3430. 3075. 3285. 3485. 3130. 1918. 2128. 2535. 2818. 2778. 2756. 2800.
3366. 2579. 2658. 2707. 2758. 2808. 2847. 2050. 2120. 2240. 2145. 2190.
2340. 2510. 2290. 2455. 2420. 1985. 2015. 2280. 2081. 2109. 2275. 2094.
2122. 2140. 2169. 2204. 2265. 2300. 2540. 2536. 2551. 2679. 2714. 2975.
2326. 2480. 2414. 2458. 2976. 3016. 3131. 2261. 2209. 2264. 2212. 2319.
2221. 2661. 2563. 2912. 2935. 3042. 3157. 2952. 3049. 3012. 3217. 3062.
2040. 2650. 2254. 1950. 2460. 3045. 2695. 2008. 2365. 1967. 3252. 3034.
1488. 2926. 3151. 3110.]

Unique values of engine-size: [nan 130. 152. 109. 136. 131. 108. 164. 209. 61. 90. 98. 122. 156.
92. 79. 110. 111. 119. 258. 326. 91. 70. 80. 140. 134. 183. 234.
308. 304. 103. 97. 120. 181. 151. 194. 203. 132. 121. 146. 171. 161.
141. 173. 145.]

Unique values of bore: [3.47 nan 3.19 3.13 3.5 3.31 3.62 2.91 3.03 2.97 3.34 3.6 2.92 3.15
3.63 3.54 3.08 3.39 3.76 3.43 3.58 3.46 3.8 3.78 3.17 3.35 3.59 2.99
3.33 3.7 3.61 3.94 3.74 2.54 3.05 3.27 3.24 3.01 2.68]

Unique values of stroke: [2.68 3.47 nan 3.4 2.8 3.19 3.39 3.03 3.11 3.23 3.46 3.9 3.41 3.07
3.58 4.17 2.76 3.15 3.16 3.64 3.1 3.35 3.12 3.86 3.29 3.27 3.52 2.19
3.21 2.9 2.07 2.36 2.64 3.08 3.5 3.54 2.87]

Unique values of compression-ratio: [9. 10. 8. 8.5 8.3 7. 8.8 nan 9.5 9.6 9.41 9.4
7.6 9.2 10.1 9.1 8.1 8.6 22. 21.5 7.5 21.9 7.8 8.4
21. 8.7 9.31 9.3 7.7 22.5 23. 22.7 11.5]

Unique values of horsepower: [111. 154. 102. 115. 110. nan 140. 160. 101. 121. 182. 48. 70. 68.
88. 145. 76. 60. 86. 100. 78. 90. 176. 262. 135. 84. 64. 120.
72. 123. 155. 184. 175. 116. 55. 69. 97. 152. 200. 95. 143. 207.
288. 73. 82. 94. 62. 56. 112. 92. 161. 156. 85. 52. 114. 162.
134. 106. 142. 58.]

Unique values of peak-rpm: [nan 5500. 5800. 4250. 5400. 5100. 5000. 4800. 6000. 4750. 4650. 4200.
4350. 4500. 5200. 4150. 5600. 5900. 5750. 5250. 4400. 6600. 5300. 4900.]

Unique values of city-mpg: [21. 19. 24. 18. nan 17. 16. 23. 20. 15. 47. 38. 37. 31. 49. 30. 27. 25.
13. 26. 22. 14. 45. 28. 32. 35. 34. 29. 36. 33.]

Unique values of highway-mpg: [27. 26. 30. 22. nan 25. 20. 29. 28. 53. 43. 41. 38. 24. 54. 42. 34. 33.
19. 17. 31. 23. 32. 18. 16. 37. 50. 39. 36. 47. 46.]

```
Unique values of price: [13495. 16500. 13950. 17450. 15250. 17710. 18920. 23875.      nan 16430.
16925. 20970. 21105. 24565. 30760. 41315.  5151.  6295.  6575.  5572.
 6377.  6229.  6692.  7609.  8558.  8921. 12964.  6479.  6855.  5399.
 6529.  7129.  7295.  7895.  9095.  8845. 10295. 12945. 10345.  6785.
35550. 36000.  6095.  6795.  6695.  7395. 10945. 11845. 13645. 15645.
 8495. 10595. 10245. 11245. 18280. 18344. 28248. 28176. 34184. 35056.
40960. 45400. 16503.  5389.  6189.  6669.  7689.  9959. 12629. 14489.
 6989.  9279.  5499.  7099.  6649.  7349.  7299.  7799.  7499.  7999.
 8249.  8949. 13499. 17199. 19699. 18399. 13200. 12440. 16900. 16695.
17075. 16630. 17950. 18150.  7957. 12764. 22018. 32528. 34028. 37028.
 9895. 12170. 15040. 15510. 18620.  5118.  7053.  7603.  7126.  7775.
 9960.  9233. 11259.  7463. 10198.  8013. 11694.  6338.  6488.  7898.
 8778.  6938.  7788.  8358.  9258.  8058.  9298.  9538.  8449.  9639.
 9989. 11199. 11549. 17669.  8948. 10698.  9988. 10898. 11248. 16558.
15998. 15690. 15750.  7975.  7995.  8195.  9495.  9995. 11595.  9980.
13295. 13845. 12290. 12940. 15985. 16515. 18950. 16845. 22470. 22625.
14399.  6849. 21485.  9295.  7198.  7738.  5195. 11900. 36880. 14869.
13860. 18420. 19045. 32250.  8499.  6918. 31600.  5348.  9549.  8238.
25552. 15580. 13415. 11048. 11850. 10795.  8189.]
```

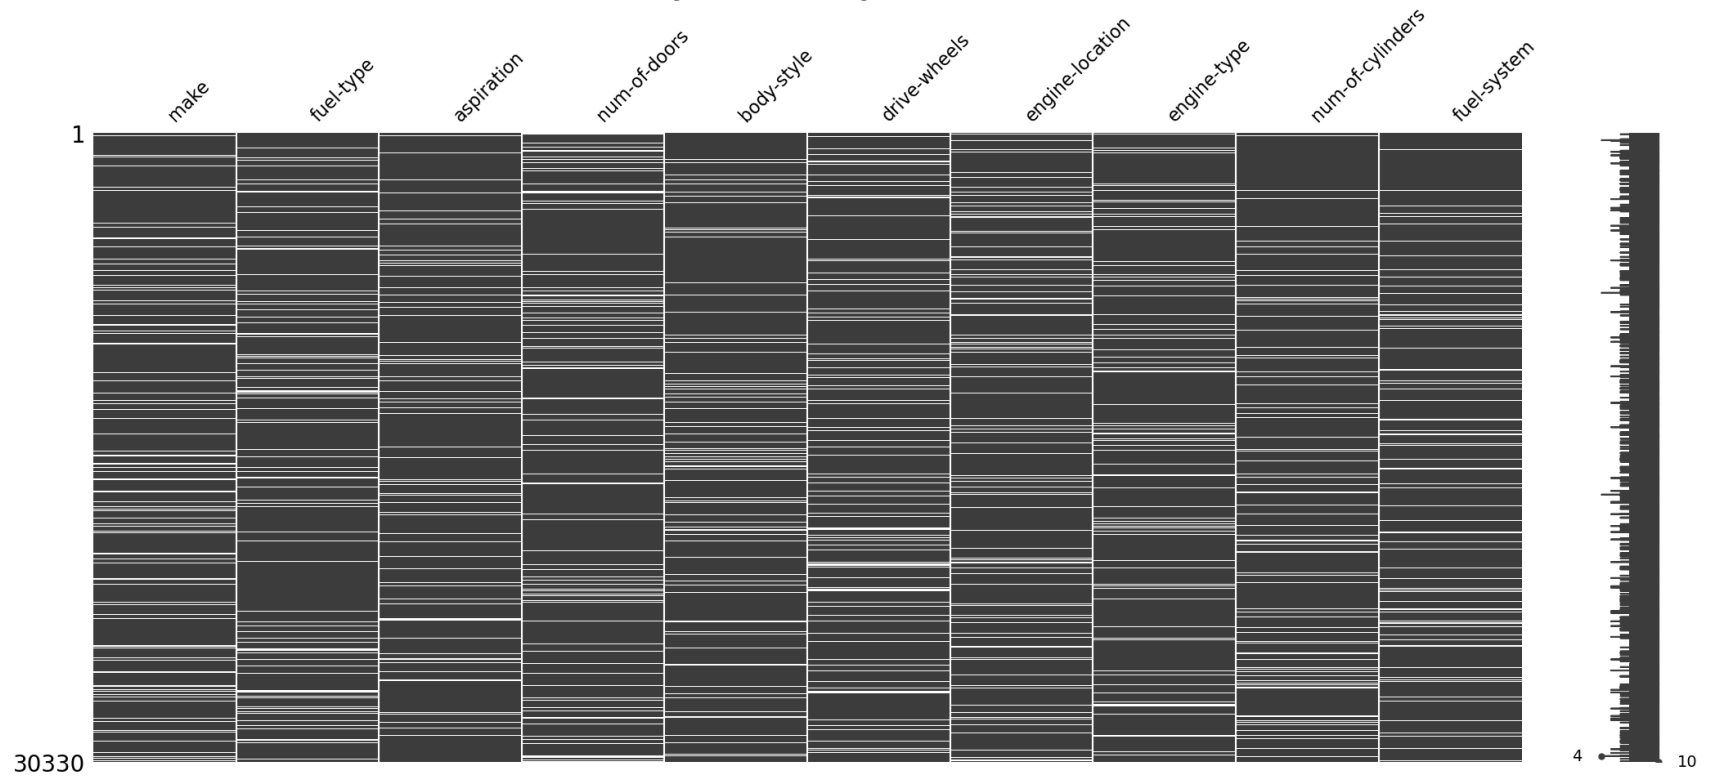
All anomalies resolved

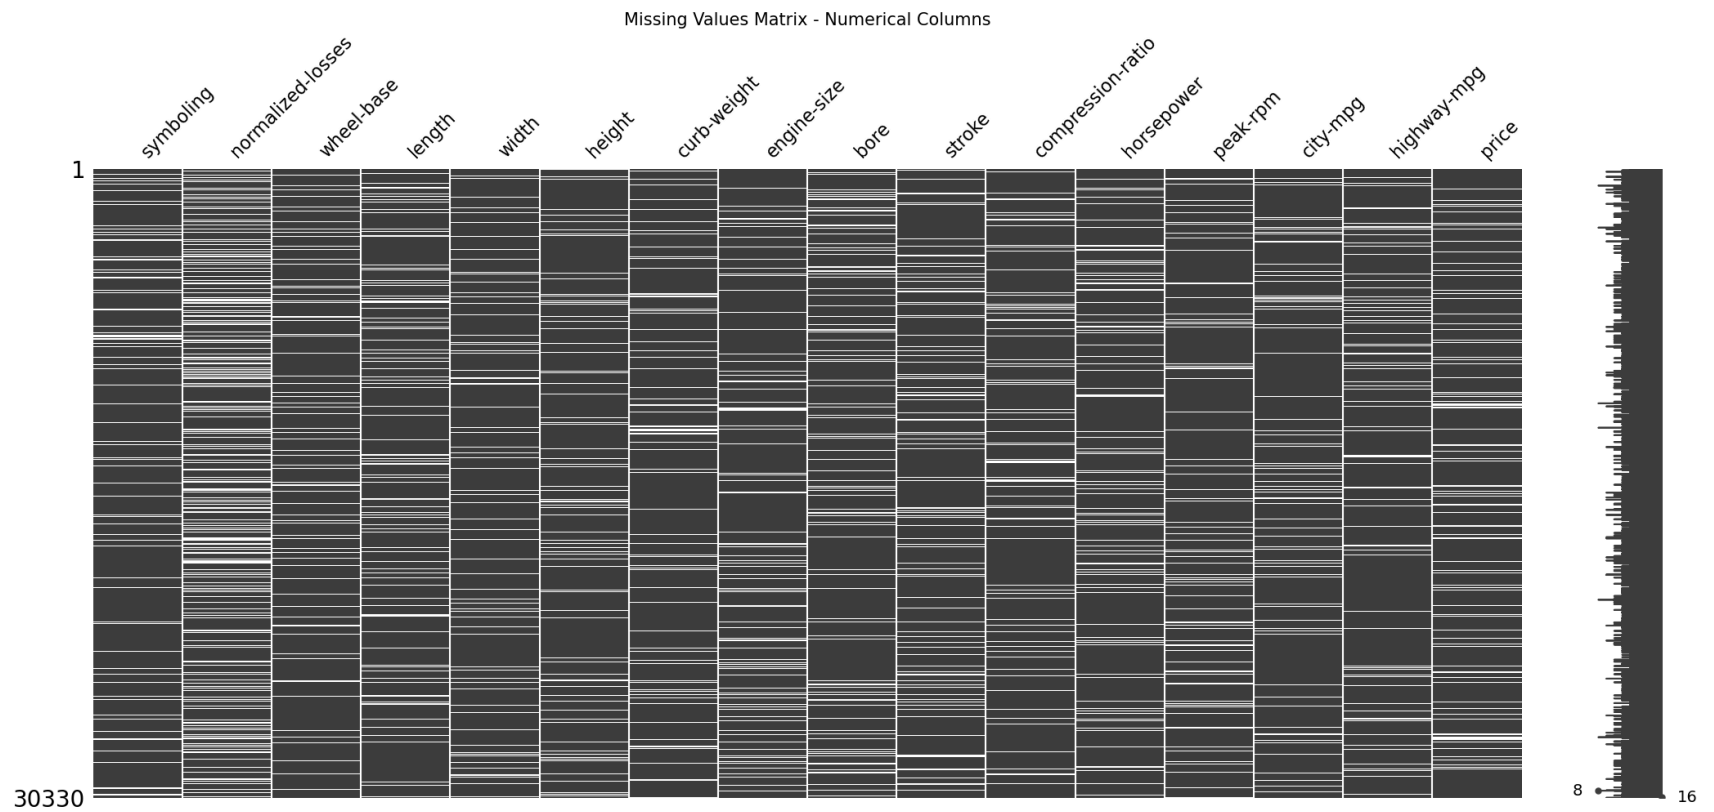
Handle Missing Values

```
In [10]: msn.matrix(df.select_dtypes(include=['object']))
plt.title('Missing Values Matrix - Categorical Columns', fontsize=15)
plt.show()

msn.matrix(df.select_dtypes(include=['number']))
plt.title('Missing Values Matrix - Numerical Columns', fontsize=15)
plt.show()
```

Missing Values Matrix - Categorical Columns





Every columns missing values is MCAR, thats a good to go

```
In [11]: (df.isnull().sum() / len(df) * 100).sort_values(ascending=False).reset_index().rename(columns={'index': 'feature', 0:
```

Out[11]:

| | feature | null_percentage |
|----|-------------------|-----------------|
| 0 | normalized-losses | 28.229476 |
| 1 | price | 11.889219 |
| 2 | bore | 11.500165 |
| 3 | stroke | 11.384768 |
| 4 | horsepower | 11.236400 |
| 5 | num-of-doors | 10.778107 |
| 6 | peak-rpm | 10.712166 |
| 7 | drive-wheels | 10.270359 |
| 8 | make | 10.227498 |
| 9 | city-mpg | 10.224200 |
| 10 | fuel-system | 10.158259 |
| 11 | length | 10.128586 |
| 12 | wheel-base | 10.108803 |
| 13 | engine-size | 10.085724 |
| 14 | height | 10.052753 |
| 15 | engine-type | 10.039565 |
| 16 | symboling | 10.036268 |
| 17 | num-of-cylinders | 9.996703 |
| 18 | curb-weight | 9.983515 |
| 19 | highway-mpg | 9.980218 |
| 20 | fuel-type | 9.960435 |
| 21 | body-style | 9.904385 |

| | feature | null_percentage |
|----|-------------------|-----------------|
| 22 | compression-ratio | 9.901088 |
| 23 | engine-location | 9.818661 |
| 24 | aspiration | 9.808770 |
| 25 | width | 9.703264 |

- The 'price' column is our target variable, and we should not impute missing values for it. Filling these missing values could introduce bias into the model, leading to inaccurate predictions and affecting the integrity of the model's results
- Split num columns and cat columns then fill with their appropriate method
- Lastly we check the missing values > 20% how to deal them so no random noise occurs

```
In [12]: df = df.dropna(subset=['price'])
```

First we deal with the numerical features

```
In [13]: numerical_features = ['symboling', 'normalized-losses', 'wheel-base', 'length', 'width',
                              'height', 'curb-weight', 'engine-size', 'bore', 'stroke',
                              'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
                              'highway-mpg', 'price']
```

Lets make a Distribution Plot to look the columns behave

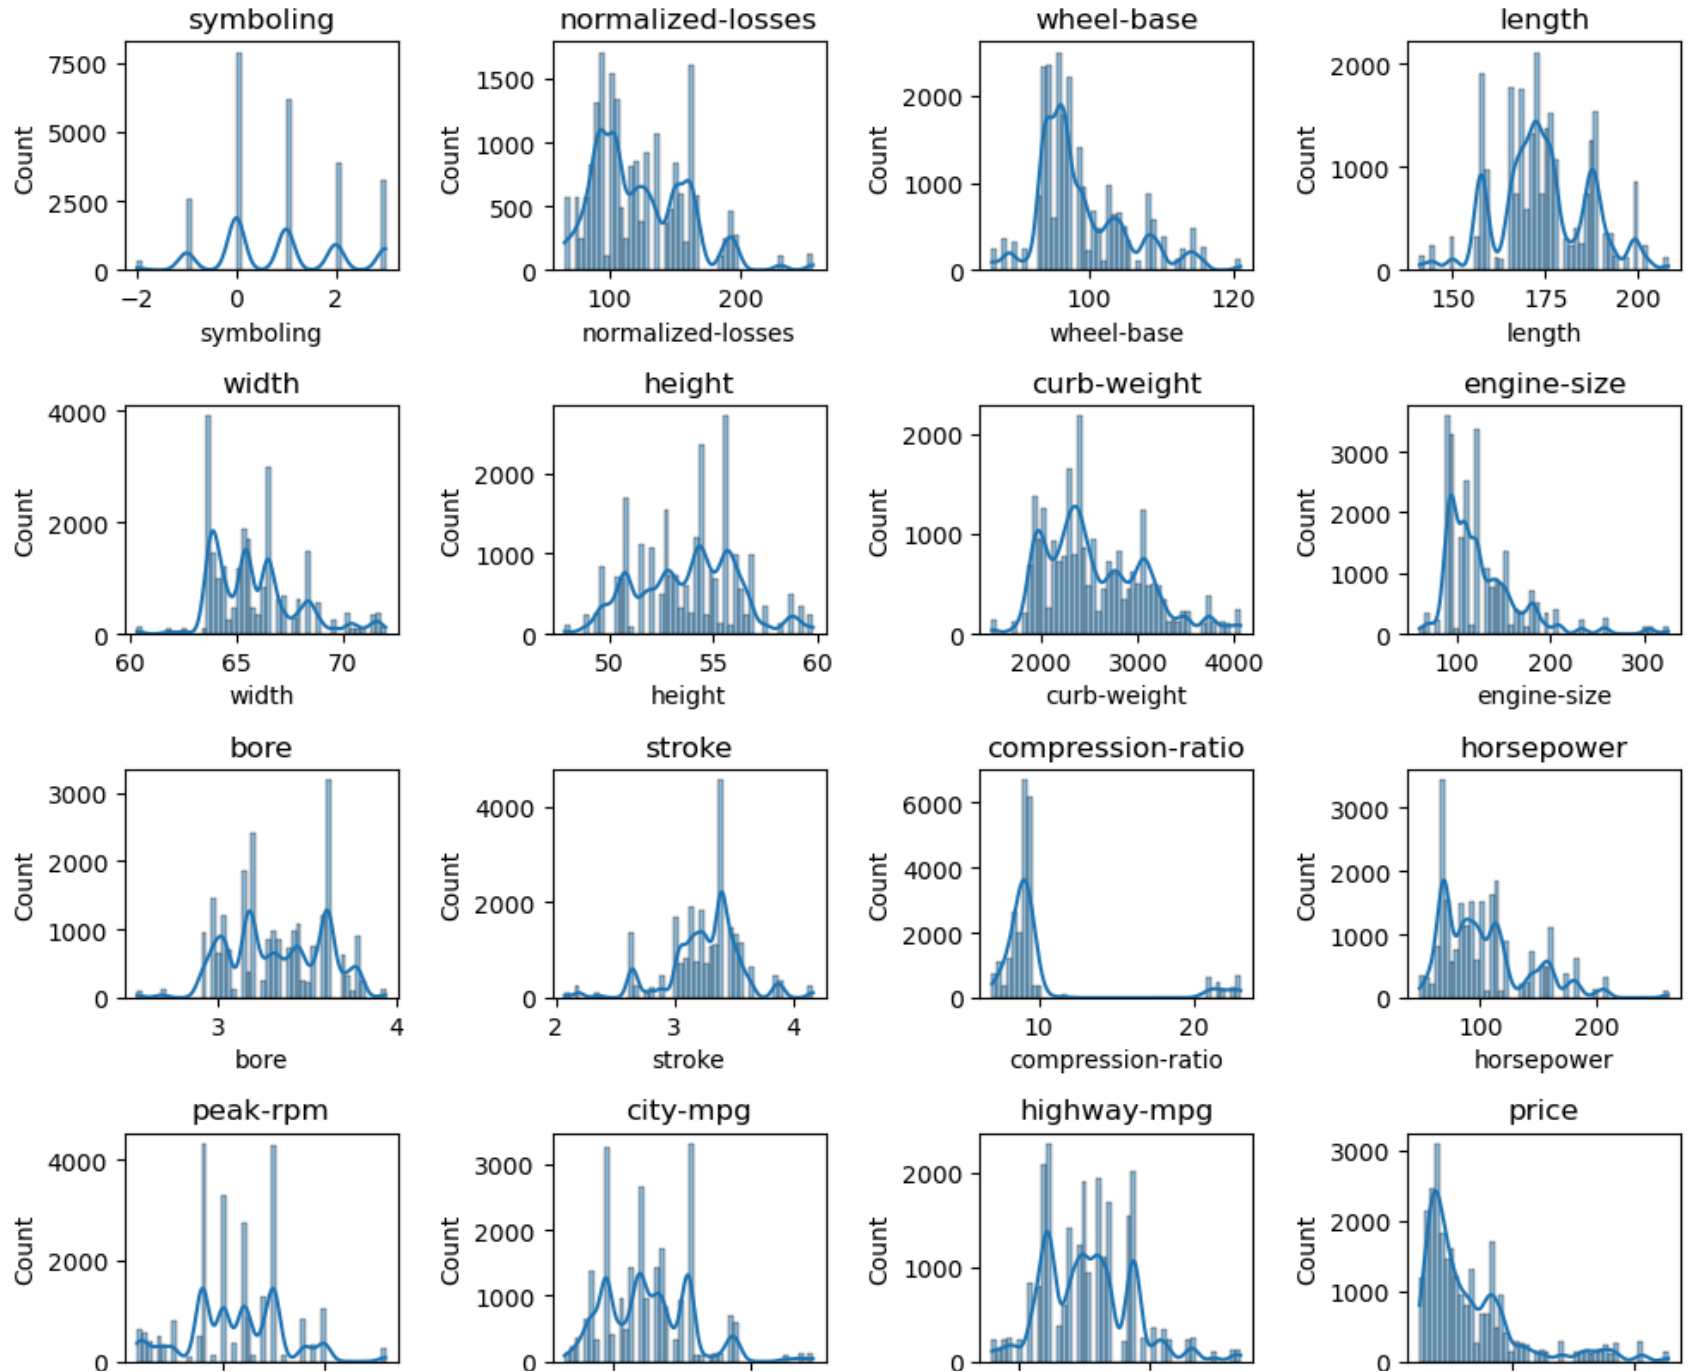
```
In [14]: plt.figure(figsize=(10, 9))

for i, col in enumerate(numerical_features):
    plt.subplot(4, 4, i+1)
    sns.histplot(df[col], bins=50, kde=True, )
    plt.title(f'{col}')

plt.suptitle('Before filling missing values', fontsize=16)

plt.tight_layout()
plt.show()
```

Before filling missing values



5000 6000
peak-rpm

20 40
city-mpg

20 40
highway-mpg

20000 40000
price

```
In [15]: without_normalzie = ['symboling', 'wheel-base', 'length', 'width',  
                             'height', 'curb-weight', 'engine-size', 'bore', 'stroke',  
                             'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',  
                             'highway-mpg']  
  
df[without_normalzie] = df[without_normalzie].fillna(df[without_normalzie].median())  
df[without_normalzie].isnull().sum()
```

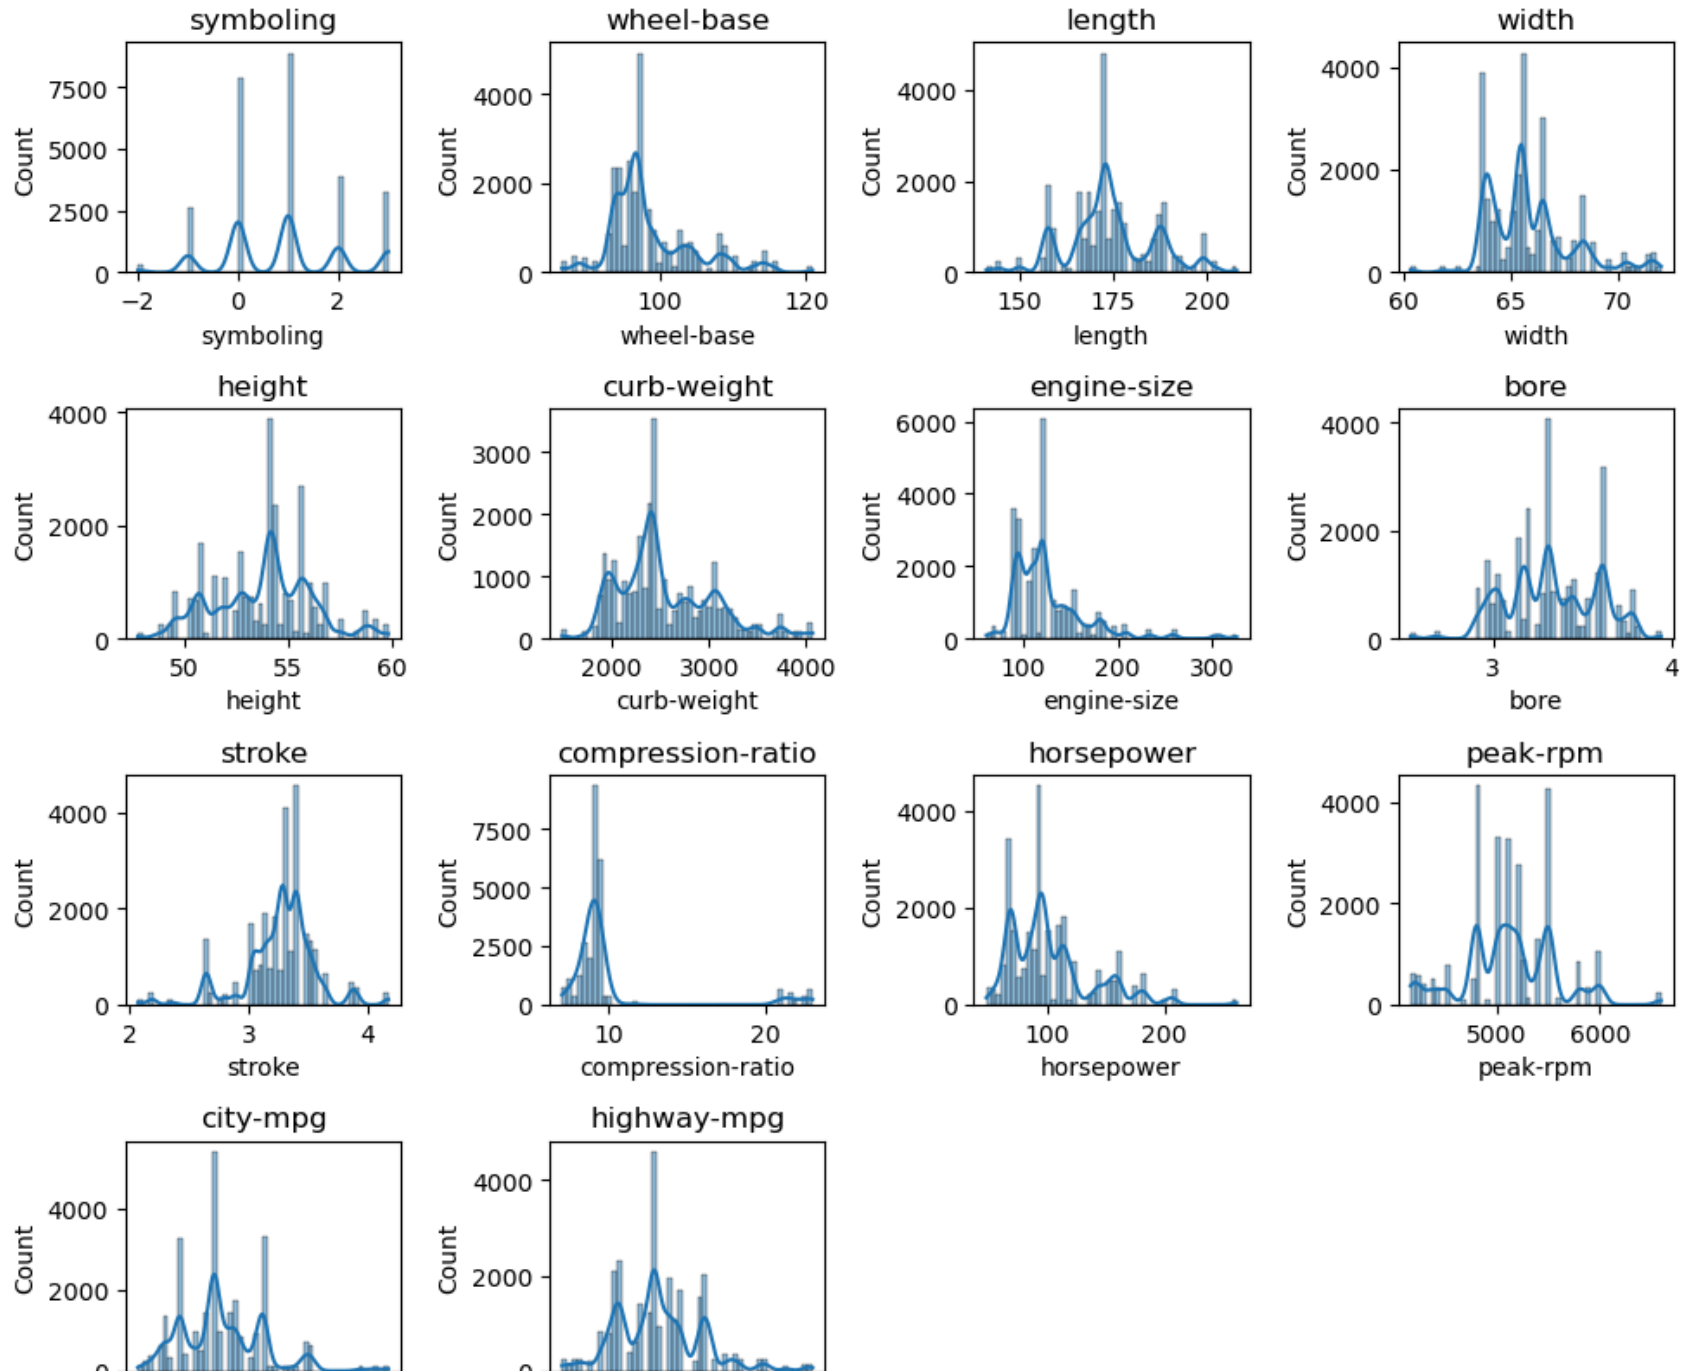
```
Out[15]: symboling      0  
wheel-base    0  
length        0  
width          0  
height         0  
curb-weight    0  
engine-size    0  
bore           0  
stroke         0  
compression-ratio 0  
horsepower     0  
peak-rpm       0  
city-mpg       0  
highway-mpg    0  
dtype: int64
```

```
In [16]: plt.figure(figsize=(10, 9))  
  
for i, col in enumerate(df[without_normalzie]):  
    plt.subplot(4, 4, i+1)  
    sns.histplot(df[col], bins=50, kde=True, )  
    plt.title(f'{col}')
```

plt.suptitle('After filling missing values', fontsize=16)

plt.tight_layout()
plt.show()

After filling missing values





Random Sampling

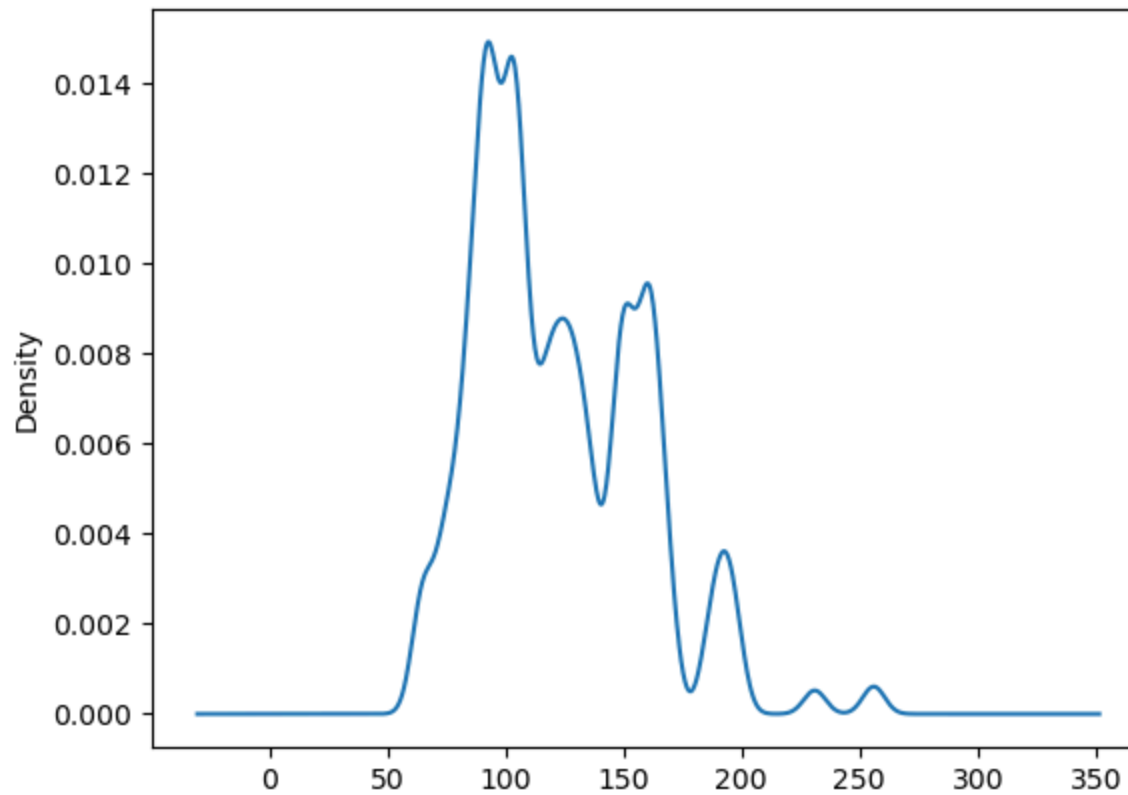
```
In [17]: missing_mask = df['normalized-losses'].isnull()

imputed_values = df.loc[~missing_mask, 'normalized-losses'].sample(
    n=missing_mask.sum(), replace=True, random_state=42)

df.loc[missing_mask, 'normalized-losses'] = imputed_values.values
```

```
In [18]: df['normalized-losses'].plot(kind='kde')
```

```
Out[18]: <Axes: ylabel='Density'>
```



SO the distribution remains same as after filling missing values

Handle categorical features

```
In [19]: categorical_features= ['make', 'fuel-type', 'aspiration', 'num-of-doors', 'body-style',  
                                'drive-wheels', 'engine-location', 'engine-type', 'num-of-cylinders',  
                                'fuel-system']
```

```
In [20]: cat_columns_missing_values = (df[categorical_features].isnull().sum() / len(df) * 100).sort_values(ascending=False).  
cat_columns_missing_values.columns = ['Cat Features', 'Missing Percentage %']  
cat_columns_missing_values
```

```
Out[20]:
```

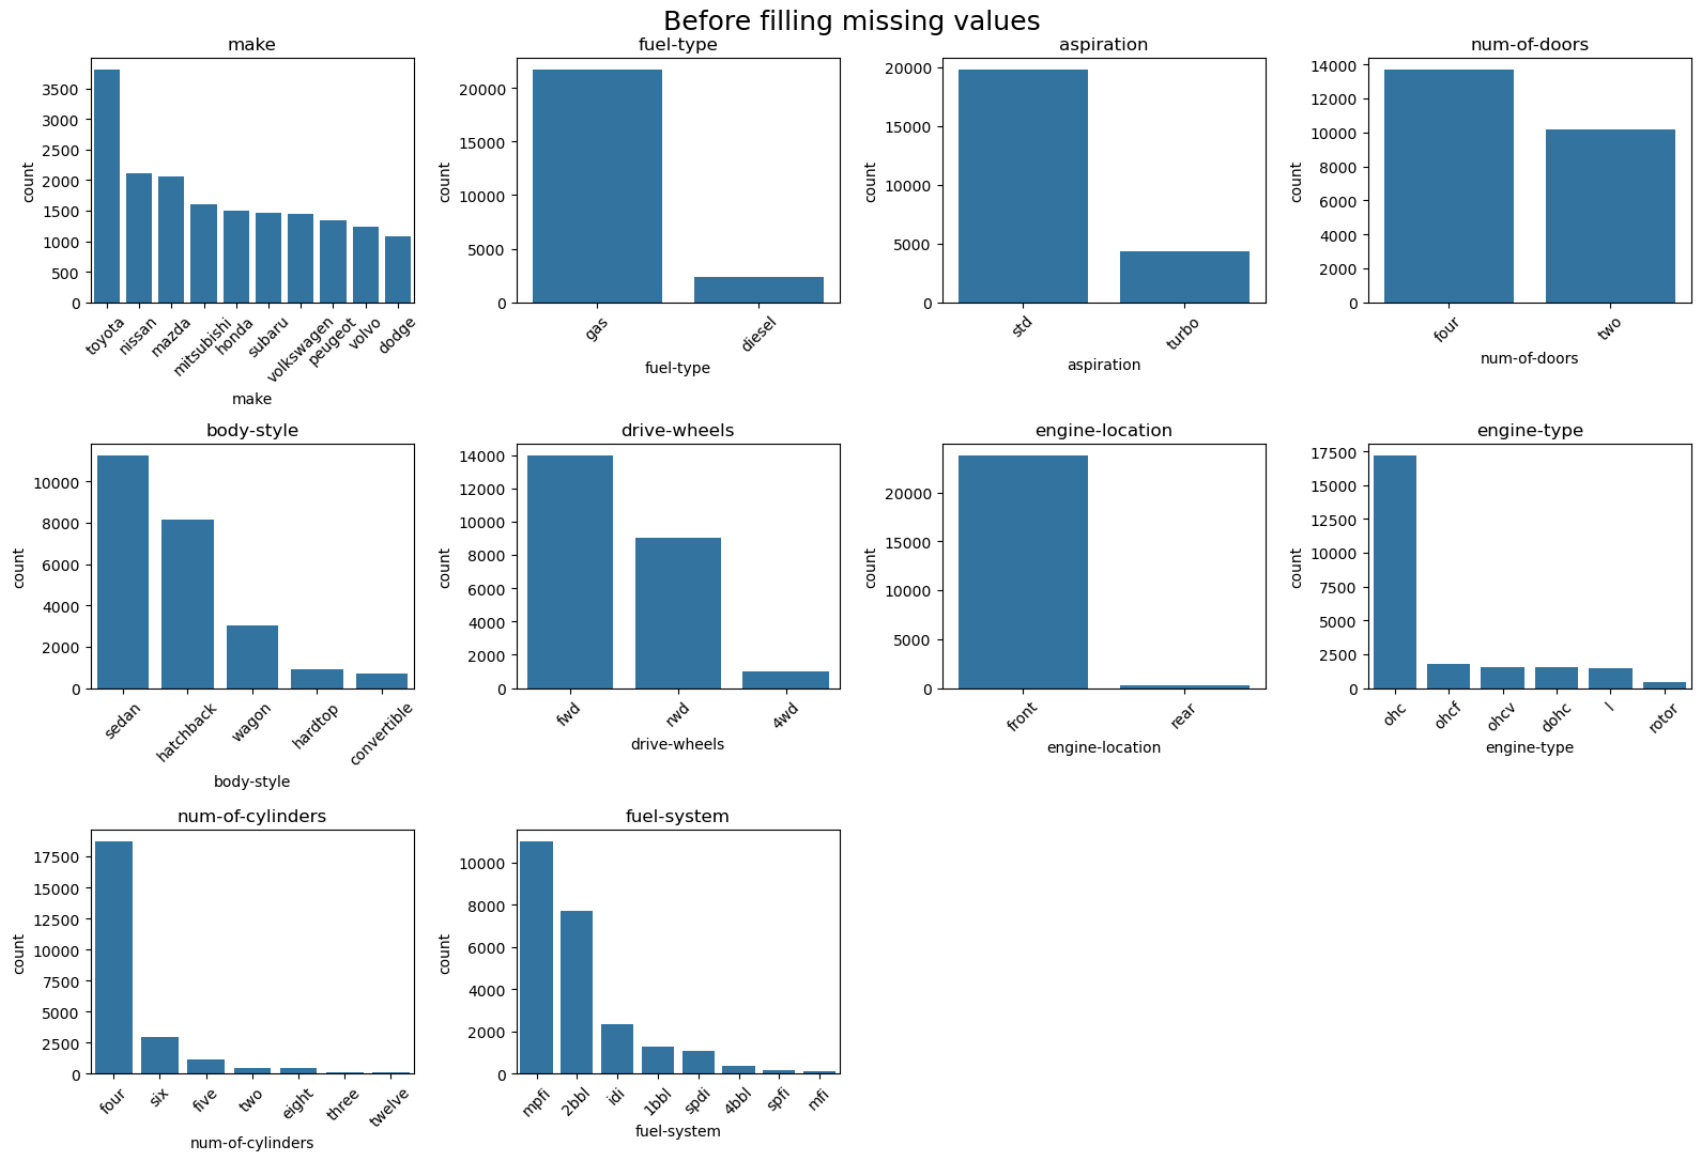
| | Cat Features | Missing Percentage % |
|---|------------------|----------------------|
| 0 | num-of-doors | 10.780572 |
| 1 | drive-wheels | 10.223021 |
| 2 | make | 10.189343 |
| 3 | fuel-system | 10.163149 |
| 4 | engine-type | 10.155665 |
| 5 | num-of-cylinders | 10.028439 |
| 6 | fuel-type | 9.991019 |
| 7 | engine-location | 9.897470 |
| 8 | body-style | 9.848825 |
| 9 | aspiration | 9.747792 |

```
In [21]: df['make'].replace('peugot', 'peugeot', inplace=True)
```

```
In [22]: plt.figure(figsize=(16, 14))  
  
for i, col in enumerate(categorical_features):  
    plt.subplot(4, 4, i+1)  
    sns.countplot(x=col, data=df, order=df[col].value_counts().head(10).index)
```

```
plt.title(f'{col}')  
plt.xticks(rotation=45)
```

```
plt.suptitle('Before filling missing values', fontsize=18)  
plt.tight_layout()  
plt.show()
```



```
In [23]: for col in categorical_features:
          df[col] = df[col].fillna(df[col].mode()[0])
```

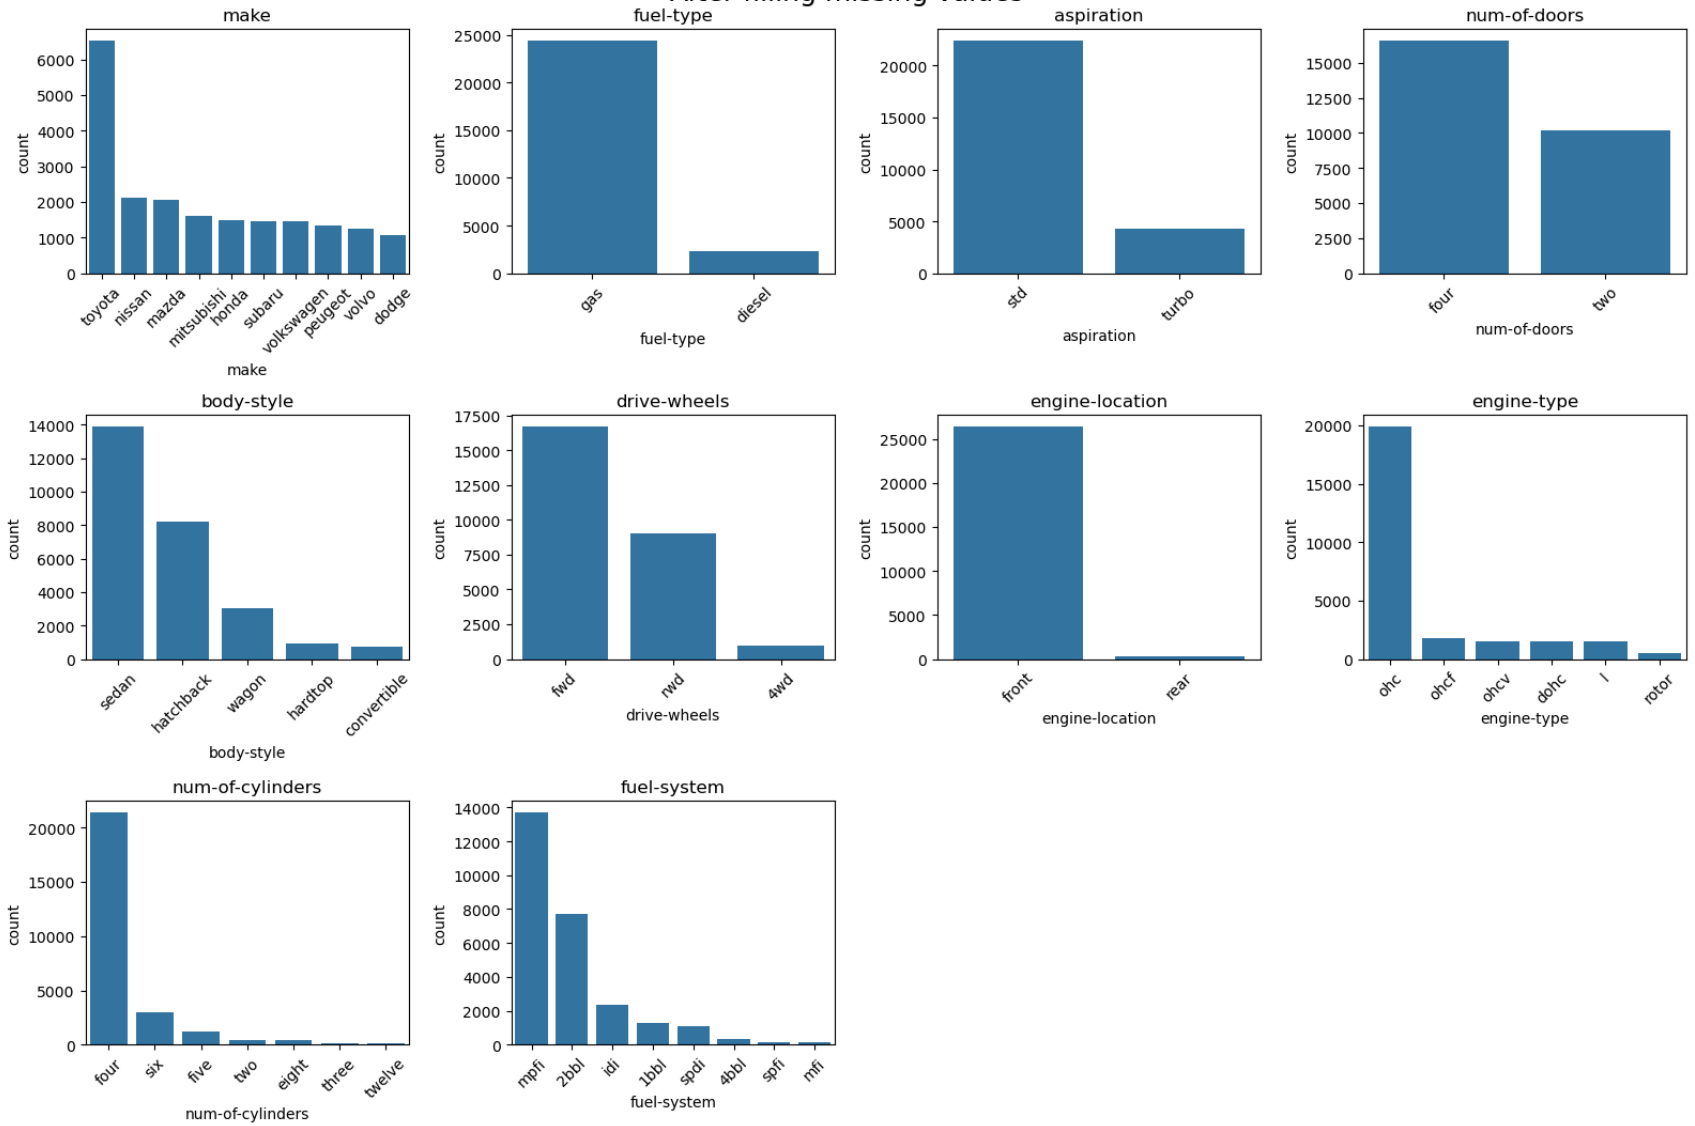
```
In [24]: plt.figure(figsize=(16, 14))

          for i, col in enumerate(df[categorical_features]):
```

```
plt.subplot(4, 4, i+1)
sns.countplot(x=col, data=df, order=df[col].value_counts().head(10).index)
plt.title(f'{col}')
plt.xticks(rotation=45)

plt.suptitle('After filling missing values', fontsize=18)
plt.tight_layout()
plt.show()
```

After filling missing values

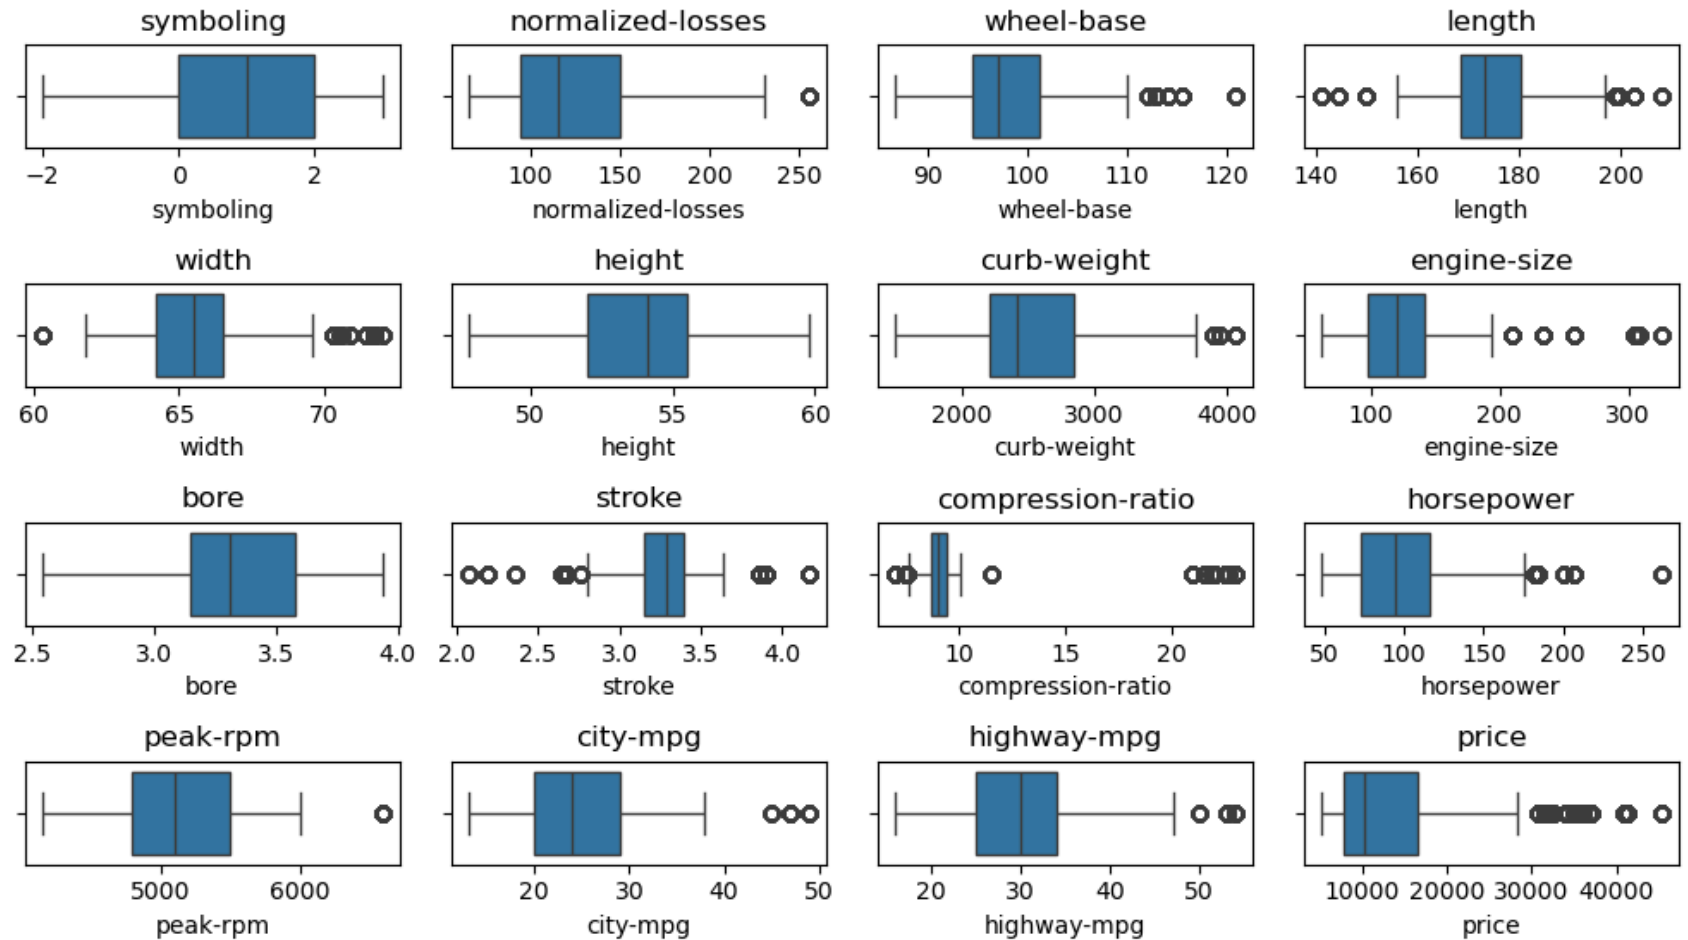


DONE

Deal With Outliers

```
In [25]: plt.figure(figsize=(10, 6))
         for i, col in enumerate(numerical_features):
             ax = plt.subplot(4, 4, i+1)
             sns.boxplot(x=df[col])
             plt.title(col)
         plt.suptitle("Boxplots of Numerical Features 'Before'")
         plt.tight_layout()
         plt.show()
```

Boxplots of Numerical Features 'Before'



```
In [26]: outliers = {}

for col in df.select_dtypes(include='number').columns:
    q1 = df[col].quantile(0.25)
    q3 = df[col].quantile(0.75)
    IQR = q3 - q1
    upper_limit = q3 + 1.5 * IQR
    lower_limit = q1 - 1.5 * IQR

    outlier_values = df[(df[col] > upper_limit) | (df[col] < lower_limit)][col]
```

```
outliers[col] = outlier_values.tolist()

print(f"\nColumn: {col}")
print(f"Upper Limit: {upper_limit}, Lower Limit: {lower_limit}")
```

Column: symboling

Upper Limit: 5.0, Lower Limit: -3.0

Column: normalized-losses

Upper Limit: 234.0, Lower Limit: 10.0

Column: wheel-base

Upper Limit: 111.25, Lower Limit: 84.44999999999999

Column: length

Upper Limit: 197.70000000000005, Lower Limit: 151.29999999999995

Column: width

Upper Limit: 69.94999999999999, Lower Limit: 60.75000000000001

Column: height

Upper Limit: 60.75, Lower Limit: 46.75

Column: curb-weight

Upper Limit: 3799.5, Lower Limit: 1259.5

Column: engine-size

Upper Limit: 205.5, Lower Limit: 33.5

Column: bore

Upper Limit: 4.225000000000005, Lower Limit: 2.505

Column: stroke

Upper Limit: 3.775, Lower Limit: 2.775

Column: compression-ratio

Upper Limit: 10.45000000000003, Lower Limit: 7.649999999999998

Column: horsepower

Upper Limit: 180.5, Lower Limit: 8.5

Column: peak-rpm

Upper Limit: 6550.0, Lower Limit: 3750.0

Column: city-mpg

Upper Limit: 42.5, Lower Limit: 6.5

Column: highway-mpg

Upper Limit: 47.5, Lower Limit: 11.5

Column: price

Upper Limit: 29587.5, Lower Limit: -5312.5

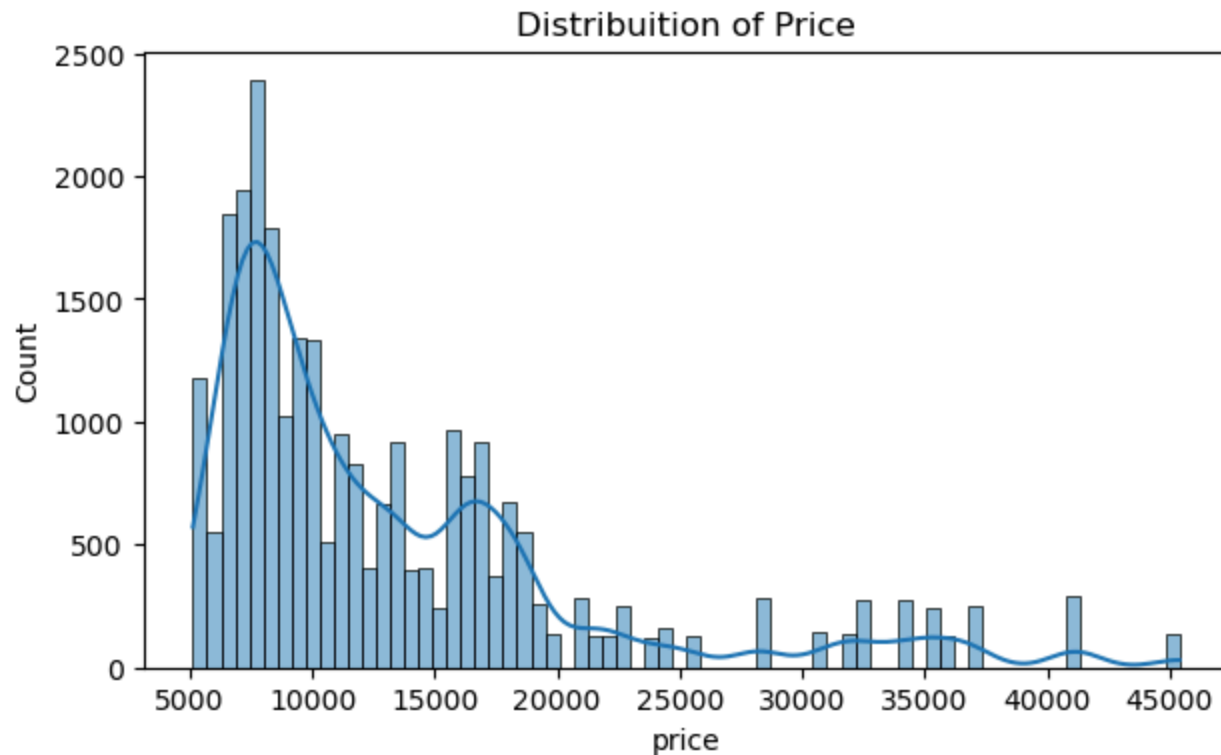
- **normalized-losses** is a numeric feature representing the relative average loss payment per insured vehicle (i.e., risk rating) for a particular car make and model, compared to other vehicles.
 - Low (e.g., 65) Lower risk of loss → lower insurance cost
 - High (e.g., 250) Higher risk of loss → higher insurance cost
- **Wheelbase** is the distance between the centers of the front and rear wheels of a vehicle.
 - Compact Car (e.g., Honda Fit): shorter wheelbase (~2300 mm)
 - Sedan (e.g., Toyota Camry): medium wheelbase (~2800 mm)
 - SUV/Truck (e.g., Ford F-150): long wheelbase (~3200+ mm)
- **Length** is the total distance from the front bumper to the rear bumper of the vehicle.
 - Compact cars ~155 – 170
 - Midsize cars ~170 – 185
 - Full-size/SUVs ~185 – 208+
- **Width** is the distance across the car from side to side, measured at its widest point, usually including mirrors unless stated otherwise.
- **Curb weight** is the total weight of a vehicle with all standard equipment, fluids (oil, coolant, etc.), and a full tank of fuel — but without any passengers or cargo.
 - 1,500 – 2,000 lbs Subcompact (e.g., small hatchbacks)
 - 2,000 – 3,000 lbs Sedans, compact SUVs
 - 3,000 – 4,500 lbs Full-size sedans, SUVs, trucks
 - 4,500 lbs Heavy-duty SUVs, luxury vehicles, pickups
- **engine size** the total volume of all the cylinders in the engine, and it's typically measured in liters (L) or cubic centimeters (cc)
 - 61 cc is much too small for typical car engines — it's in the realm of motorcycles or go-karts.
 - 326 cc is still quite small for a car, as most car engines range from 1.0L (1000 cc) to 5.0L and higher.

- **stroke** refers to the distance the piston travels inside the engine cylinder from the top dead center (TDC) to the bottom dead center (BDC)
 - range (2.07 to 4.17 inches) is within a realistic range for many vehicles.
 - **compression** ratio in cars is a key engine parameter that measures how much the air-fuel mixture is compressed inside the engine cylinder before ignition.
 - 7 might be used in older petrol engines or boosted engines where lower compression avoids knock.
 - 23 typical of diesel engines which require high compression for ignition.
 - **horsepower** it tells you how much work an engine can do over time.
 - 40 – 80 Very low, micro or city cars
 - 80 – 120 Economy cars, basic sedans
 - 120 – 180 Mid-range, family cars
 - 180 – 250 Sporty or performance cars
 - 250+ High-performance/sports cars
 - **peak-rpm** it measures how many times the engine's crankshaft spins in one minute.
 - 4,150 (min) Likely the peak torque RPM of a diesel engine or a low-revving petrol engine — common in utility or economy cars.
 - 6,600 (max) Likely the peak horsepower RPM of a high-revving petrol engine — could be a sporty or performance-oriented car.
 - **city-mpg** is the number of miles that a car can travel on a gallon of gasoline in the city.
 - Very low fuel efficiency, likely a large SUV, truck, or sports car with a big engine.
 - Very high fuel efficiency, likely a hybrid, compact car, or small diesel engine. Could also be an electric car equivalent, depending on dataset.
 - **highway-mpg** is the number of miles that a car can travel on a gallon of gasoline on the highway.
 - Very low fuel efficiency, likely a large SUV, truck, or sports car with a big engine.
 - Very high fuel efficiency, likely a hybrid, compact car, or small diesel engine. Could also be an electric car equivalent.
 - **price** is the price of the car in thousands of dollars.
 - That is not outliers we have some luxury cars like (mercedes, audi, bmw, etc) that price are more than 300K / 400k.
-

Exploratory Data Analysis

Univariate Analysis

```
In [27]: plt.figure(figsize=(7,4))
sns.histplot(x= df['price'], kde=True)
plt.title('Distribution of Price')
plt.show()
```



Price distribution is right skewed

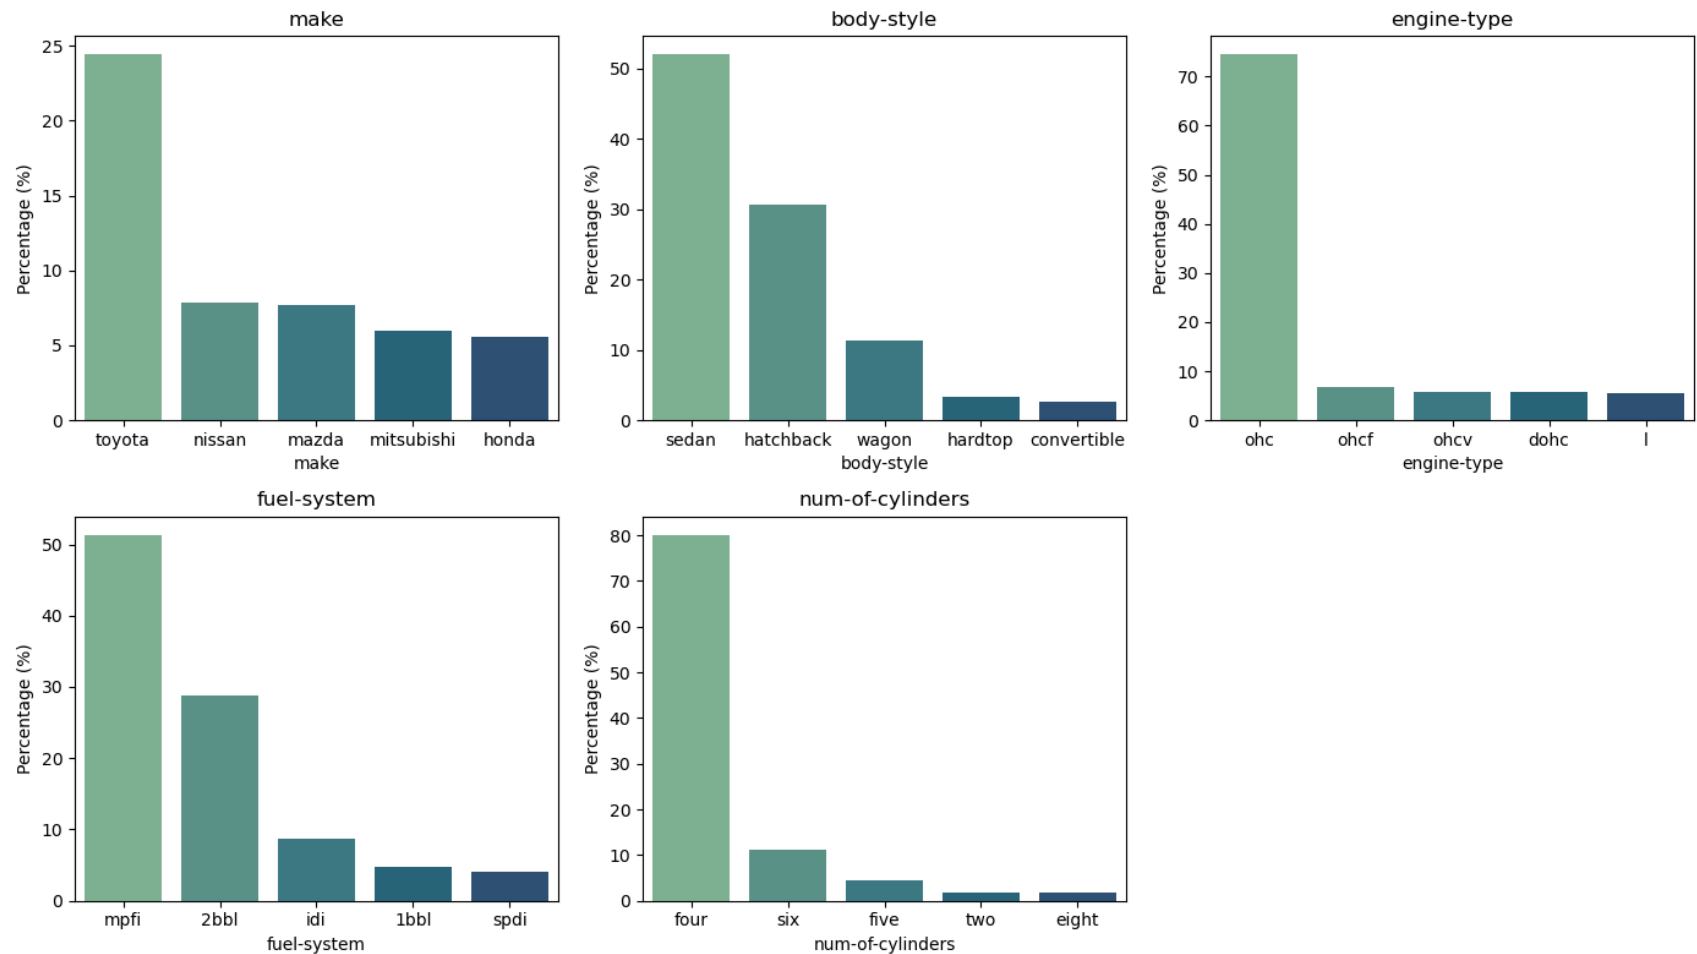
```
In [28]: a = ['make', 'body-style', 'engine-type', 'fuel-system', 'num-of-cylinders']

plt.figure(figsize=(14, 8))
for i, col in enumerate(a):
    plt.subplot(2, 3, i+1)

    category_percentage = df[col].value_counts(normalize=True).mul(100).nlargest(5)
```

```
sns.barplot(x=category_percentage.index, y=category_percentage.values, palette='crest')
plt.title(f'{col}')
plt.ylabel('Percentage (%)')
```

```
plt.tight_layout()
plt.show()
```



- Toyota emerges as the most popular car brand, comprising approximately 24% of all entries in the dataset.
- This dominance suggest Toyota's strong presence in the market or a higher availability of Toyota cars in the dataset.
- The sedan is the most common car type, making up around 50% of the dataset.

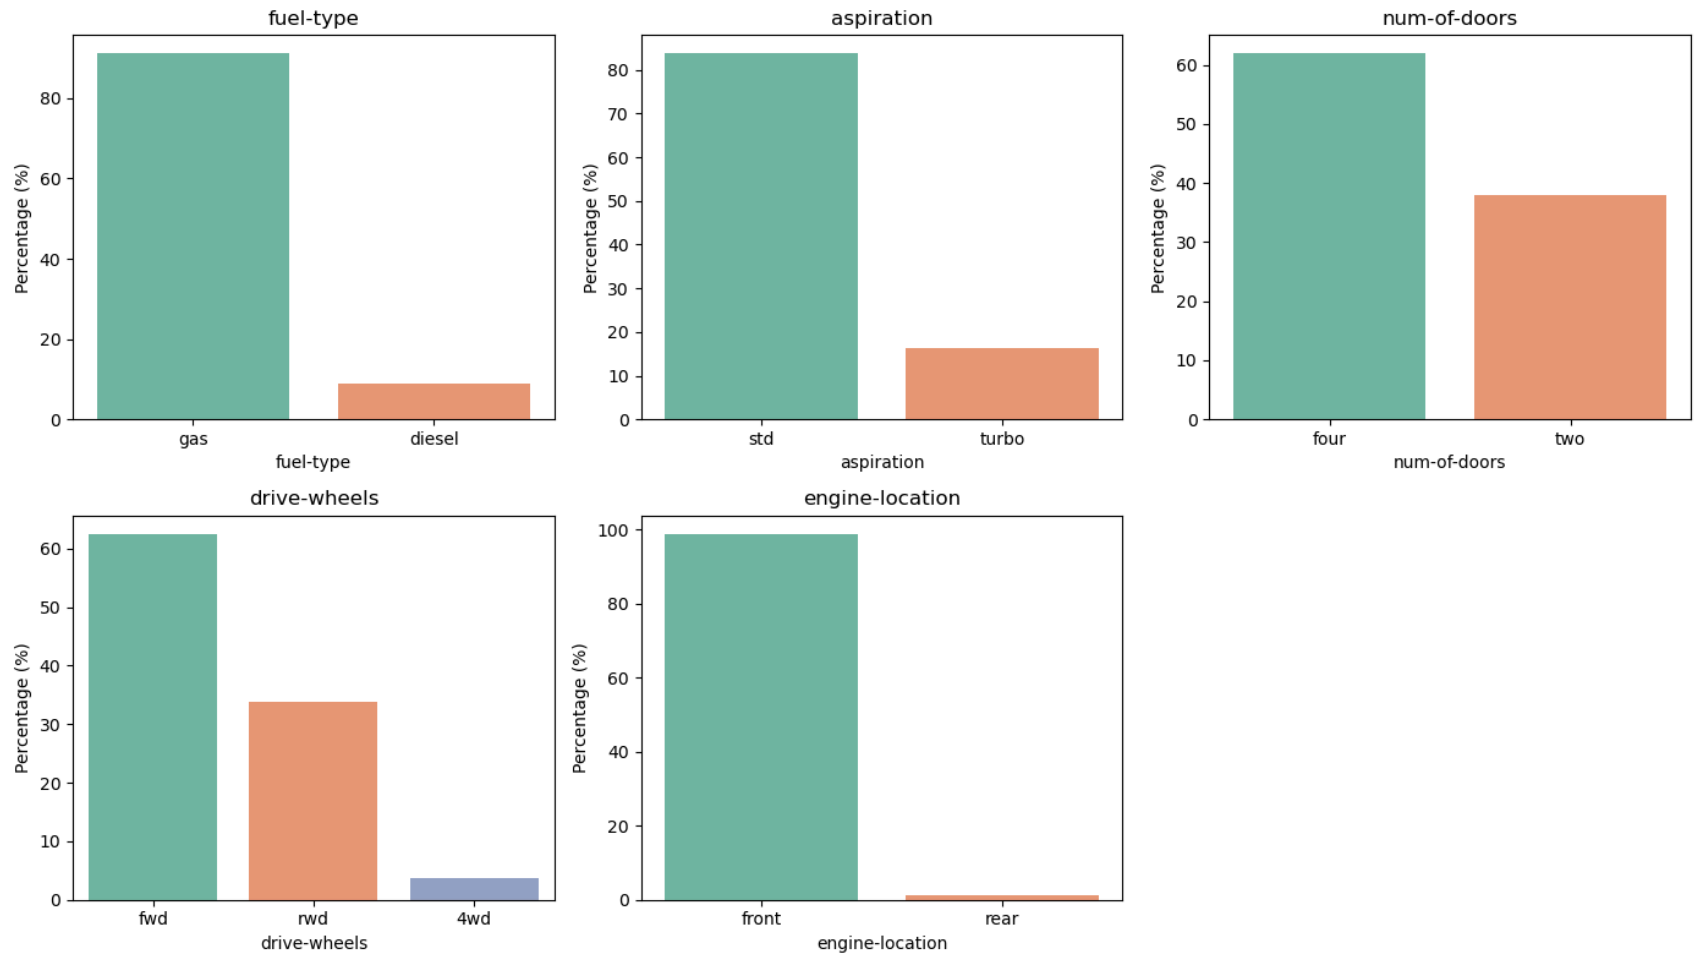
- Sedans are traditionally favored for their balance of comfort, space, and fuel efficiency, making this finding consistent with market trends.
 - An impressive 70% of the cars in the dataset feature an OHC (Overhead Camshaft) engine.
 - OHC engines are known for their reliability and efficiency, which may explain their prevalence in the dataset.
 - The MPFI (Multi-Point Fuel Injection) system is used in 50% of the vehicles in the dataset.
 - MPFI is known for improving fuel efficiency and reducing emissions, which may explain its popularity in more recent vehicle models.
 - The most common engine configuration in the dataset is the 4-cylinder engine, which accounts for a dominant 78% of the entries.
 - This high percentage suggests that 4-cylinder engines are widely preferred in the dataset, likely due to their balance of performance, fuel efficiency, and cost-effectiveness, which makes them popular in economy and compact cars.
-

```
In [29]: b = ['fuel-type', 'aspiration', 'num-of-doors', 'drive-wheels', 'engine-location']

plt.figure(figsize=(14, 8))
for i, col in enumerate(b):
    plt.subplot(2, 3, i+1)

    category_percentage = df[col].value_counts(normalize=True).mul(100).nlargest(5)
    sns.barplot(x=category_percentage.index, y=category_percentage.values, palette='Set2')
    plt.title(f'{col}')
    plt.ylabel('Percentage (%)')

plt.tight_layout()
plt.show()
```



- 90% of the cars in the dataset use gasoline as their fuel type.
- This highlights the dominance of gas-powered vehicles in the dataset, aligning with general market trends where gasoline is the most common fuel type.
- A substantial 85% of the vehicles in the dataset are equipped with standard aspiration (non-turbocharged engines).
- This indicates that turbocharged engines are relatively rare in this dataset compared to more conventional, naturally aspirated engines.
- The majority of cars, 60%, feature four doors.
- Four-door vehicles are popular due to their practicality, especially in sedans and family cars.
- 60% of the cars in the dataset have front-wheel drive (FWD).

- **Front-wheel drive** is common in compact and mid-size sedans due to its lower cost and better fuel efficiency in most driving conditions.
 - An overwhelming **95%** of the vehicles in the dataset have their engine located at the front.
-

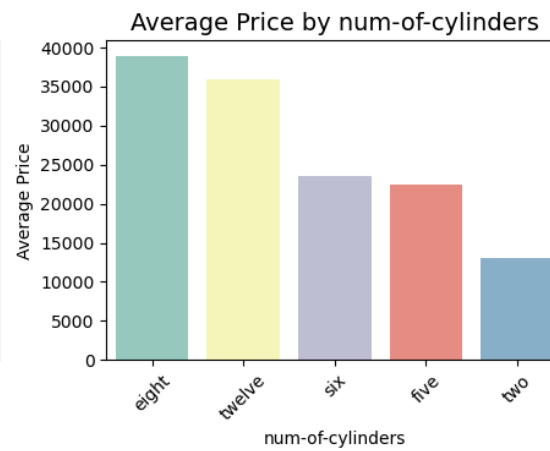
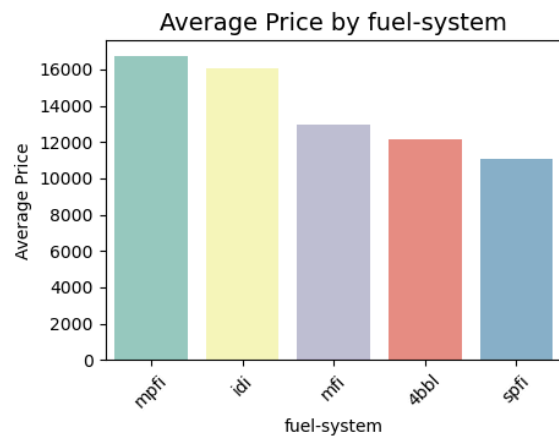
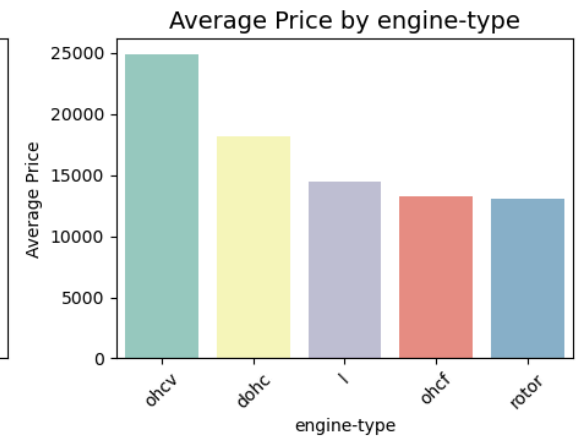
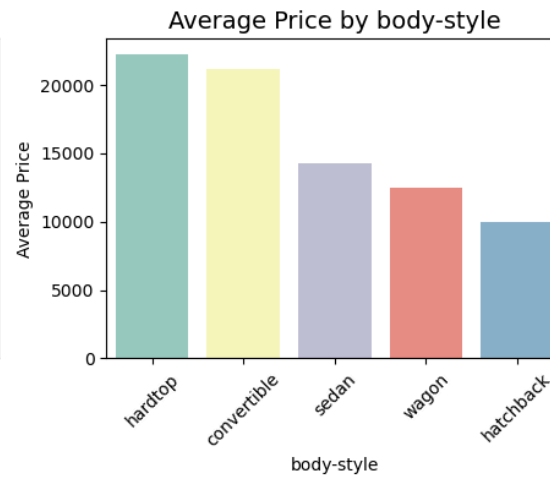
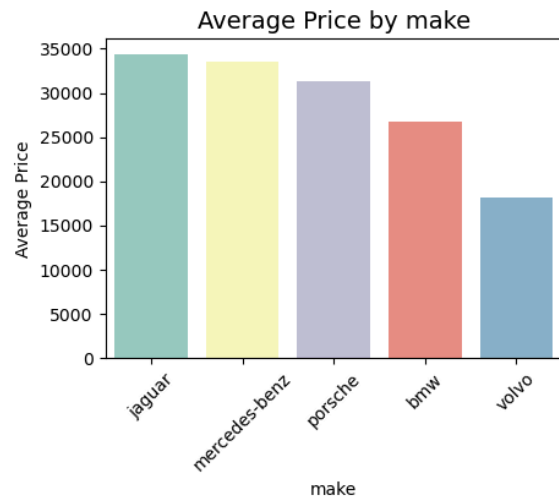
Bivariate Analysis

```
In [30]: plt.figure(figsize=(14, 8))

for i, col in enumerate(a):
    plt.subplot(2, 3, i+1)

    price_analysis = df.groupby(col)['price'].mean().reset_index().sort_values(by='price', ascending=False).head(5)
    sns.barplot(x=price_analysis[col], y=price_analysis['price'], palette='Set3')
    plt.title(f'Average Price by {col}', fontsize=14)
    plt.ylabel('Average Price')
    plt.xlabel(col)
    plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



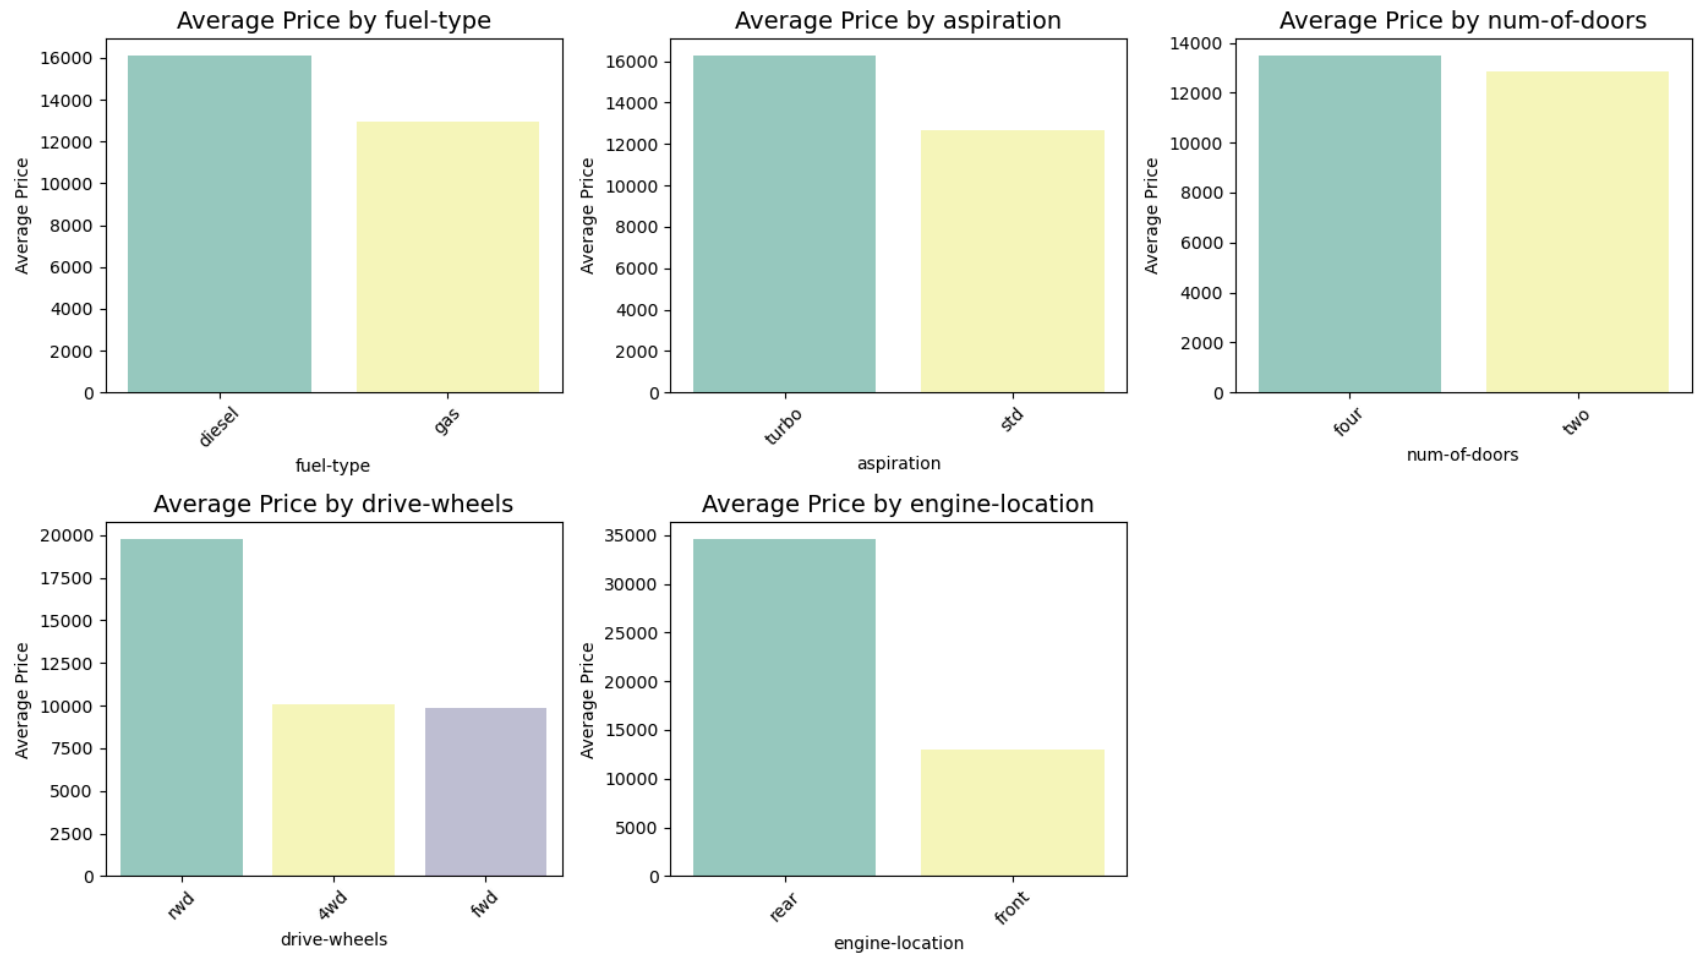
- Jaguar, Mercedes-Benz, Porsche, and BMW are the most expensive car brands in the dataset, with an average price significantly higher than others.
 - Hardtop and Convertible body styles are the most expensive, typically associated with more luxurious or performance-oriented vehicles.
 - OHCV engine type is the most expensive, likely indicating a higher-performance engine configuration.
 - MPFI (Multi-Point Fuel Injection) and IDI (Indirect Injection) fuel systems are associated with the highest prices in the dataset, suggesting these systems are typically found in higher-end vehicles.
 - Cars with 8 and 12 cylinders tend to have the highest average prices, reflecting their association with powerful, high-performance vehicles.
-

```
In [31]: plt.figure(figsize=(14, 8))

for i, col in enumerate(b):
    plt.subplot(2, 3, i+1)

    price_analysis = df.groupby(col)['price'].mean().reset_index().sort_values(by='price', ascending=False).head(5)
    sns.barplot(x=price_analysis[col], y=price_analysis['price'], palette='Set3')
    plt.title(f'Average Price by {col}', fontsize=14)
    plt.ylabel('Average Price')
    plt.xlabel(col)
    plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```



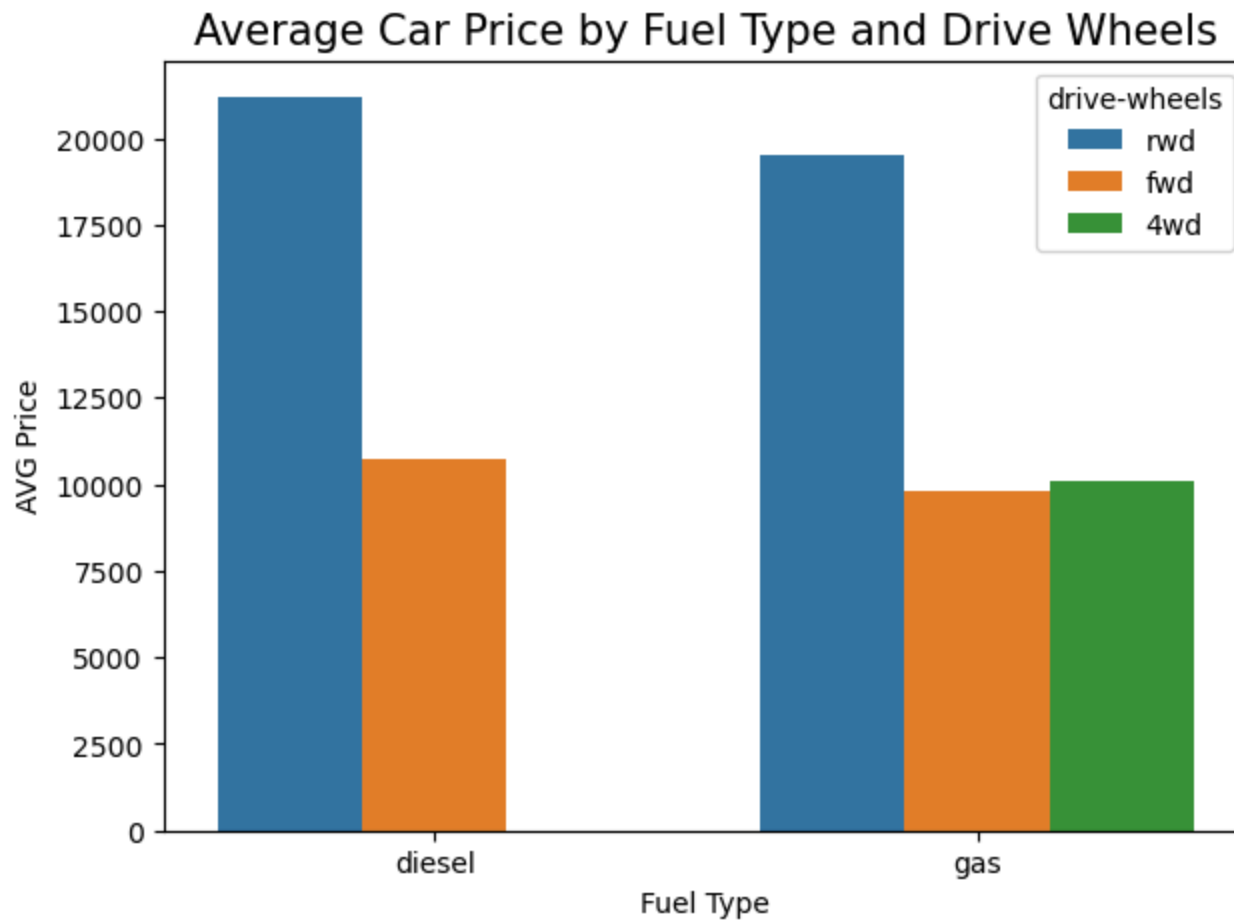
- Diesel cars tend to have a higher average price than petrol cars, likely due to their more complex engine configurations and higher efficiency for long-distance driving.
- Cars with turbocharged engines are more expensive than those with standard (naturally aspirated) engines, as turbocharged engines often provide higher performance and are used in more premium models.
- Cars with two and four doors tend to have a similar price, suggesting that the number of doors doesn't significantly impact the price in your dataset.
- RWD cars generally have a higher price than FWD cars, which could be due to their association with sports and luxury cars that typically feature RWD.

- Cars with a rear engine location are more expensive than those with a front engine location, likely due to the unique engineering and performance benefits of rear-engine cars, such as those found in high-end sports cars.
-

Multivariate Analysis

```
In [32]: fuel_drive_price = df.groupby(['fuel-type', 'drive-wheels'])['price'].mean().reset_index().sort_values(by='price', ascending=True)

plt.figure(figsize=(7,5))
sns.barplot(x='fuel-type', y='price', data=fuel_drive_price, hue='drive-wheels')
plt.title('Average Car Price by Fuel Type and Drive Wheels', fontsize=15)
plt.xlabel('Fuel Type')
plt.ylabel('AVG Price')
plt.show()
```



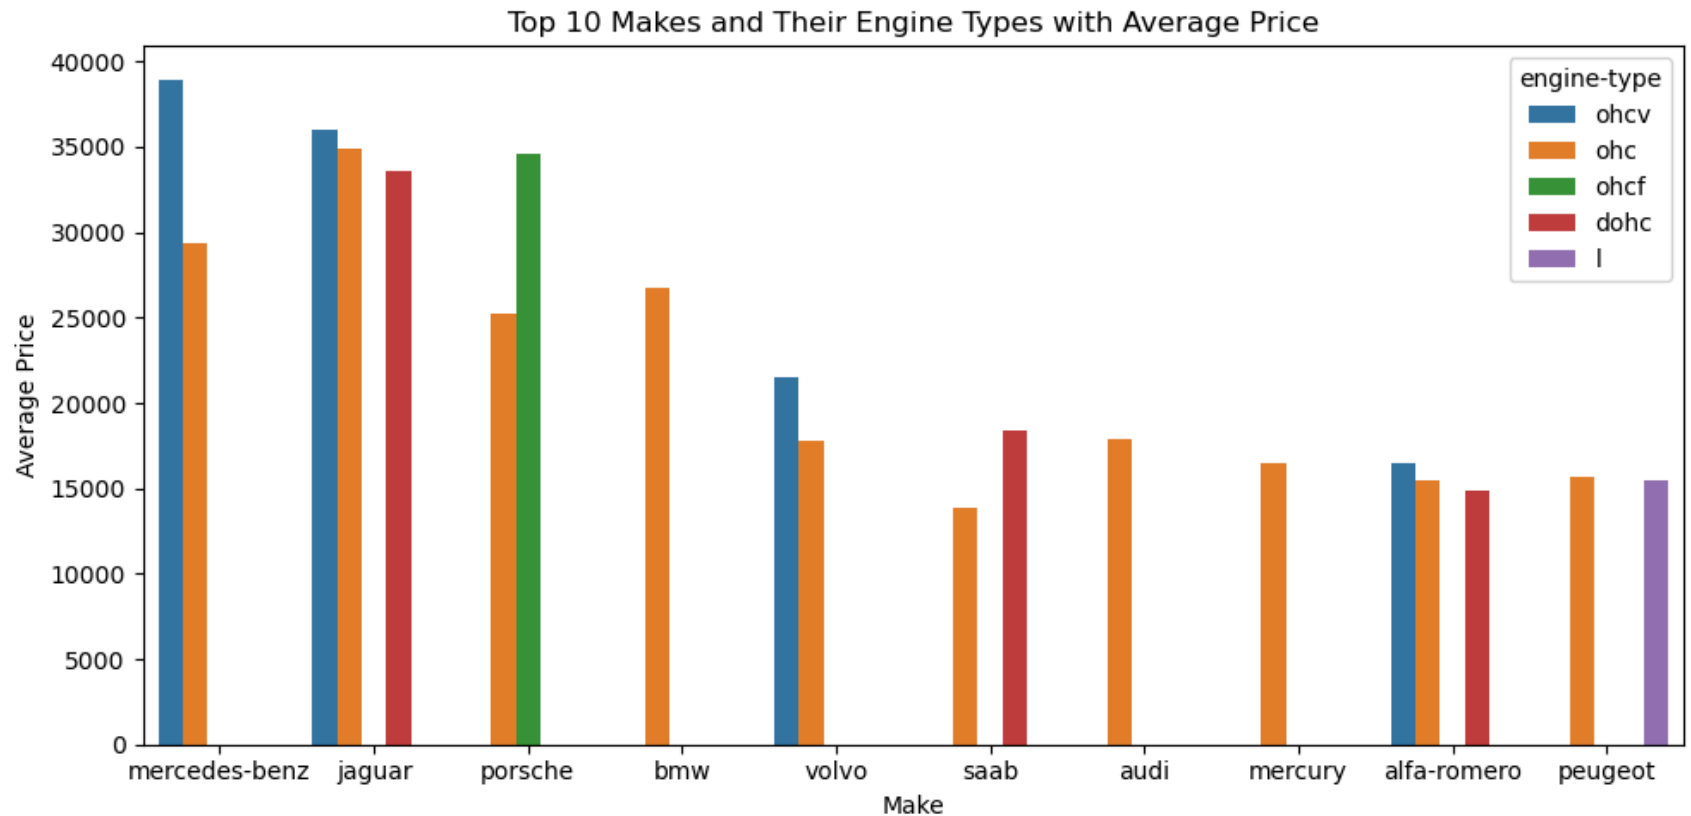
Rear-wheel drive (RWD) cars tend to be more expensive, and most of them are available in both diesel and gas fuel types. On the other hand, four-wheel drive (4WD) cars are only available with gas engines.

```
In [33]: make_engine_avgprice = df.groupby(['make', 'engine-type'])['price'].mean().sort_values(ascending=False).reset_index()
top_10_makes = make_engine_avgprice.groupby('make')['price'].mean().sort_values(ascending=False).head(10).index
top_10_make_engine_avgprice = make_engine_avgprice[make_engine_avgprice['make'].isin(top_10_makes)]

plt.figure(figsize=(10, 5))
sns.barplot(x='make', y='price', data=top_10_make_engine_avgprice, hue='engine-type')
plt.title('Top 10 Makes and Their Engine Types with Average Price')
```



```
plt.xlabel('Make')
plt.ylabel('Average Price')
plt.tight_layout()
plt.show()
```



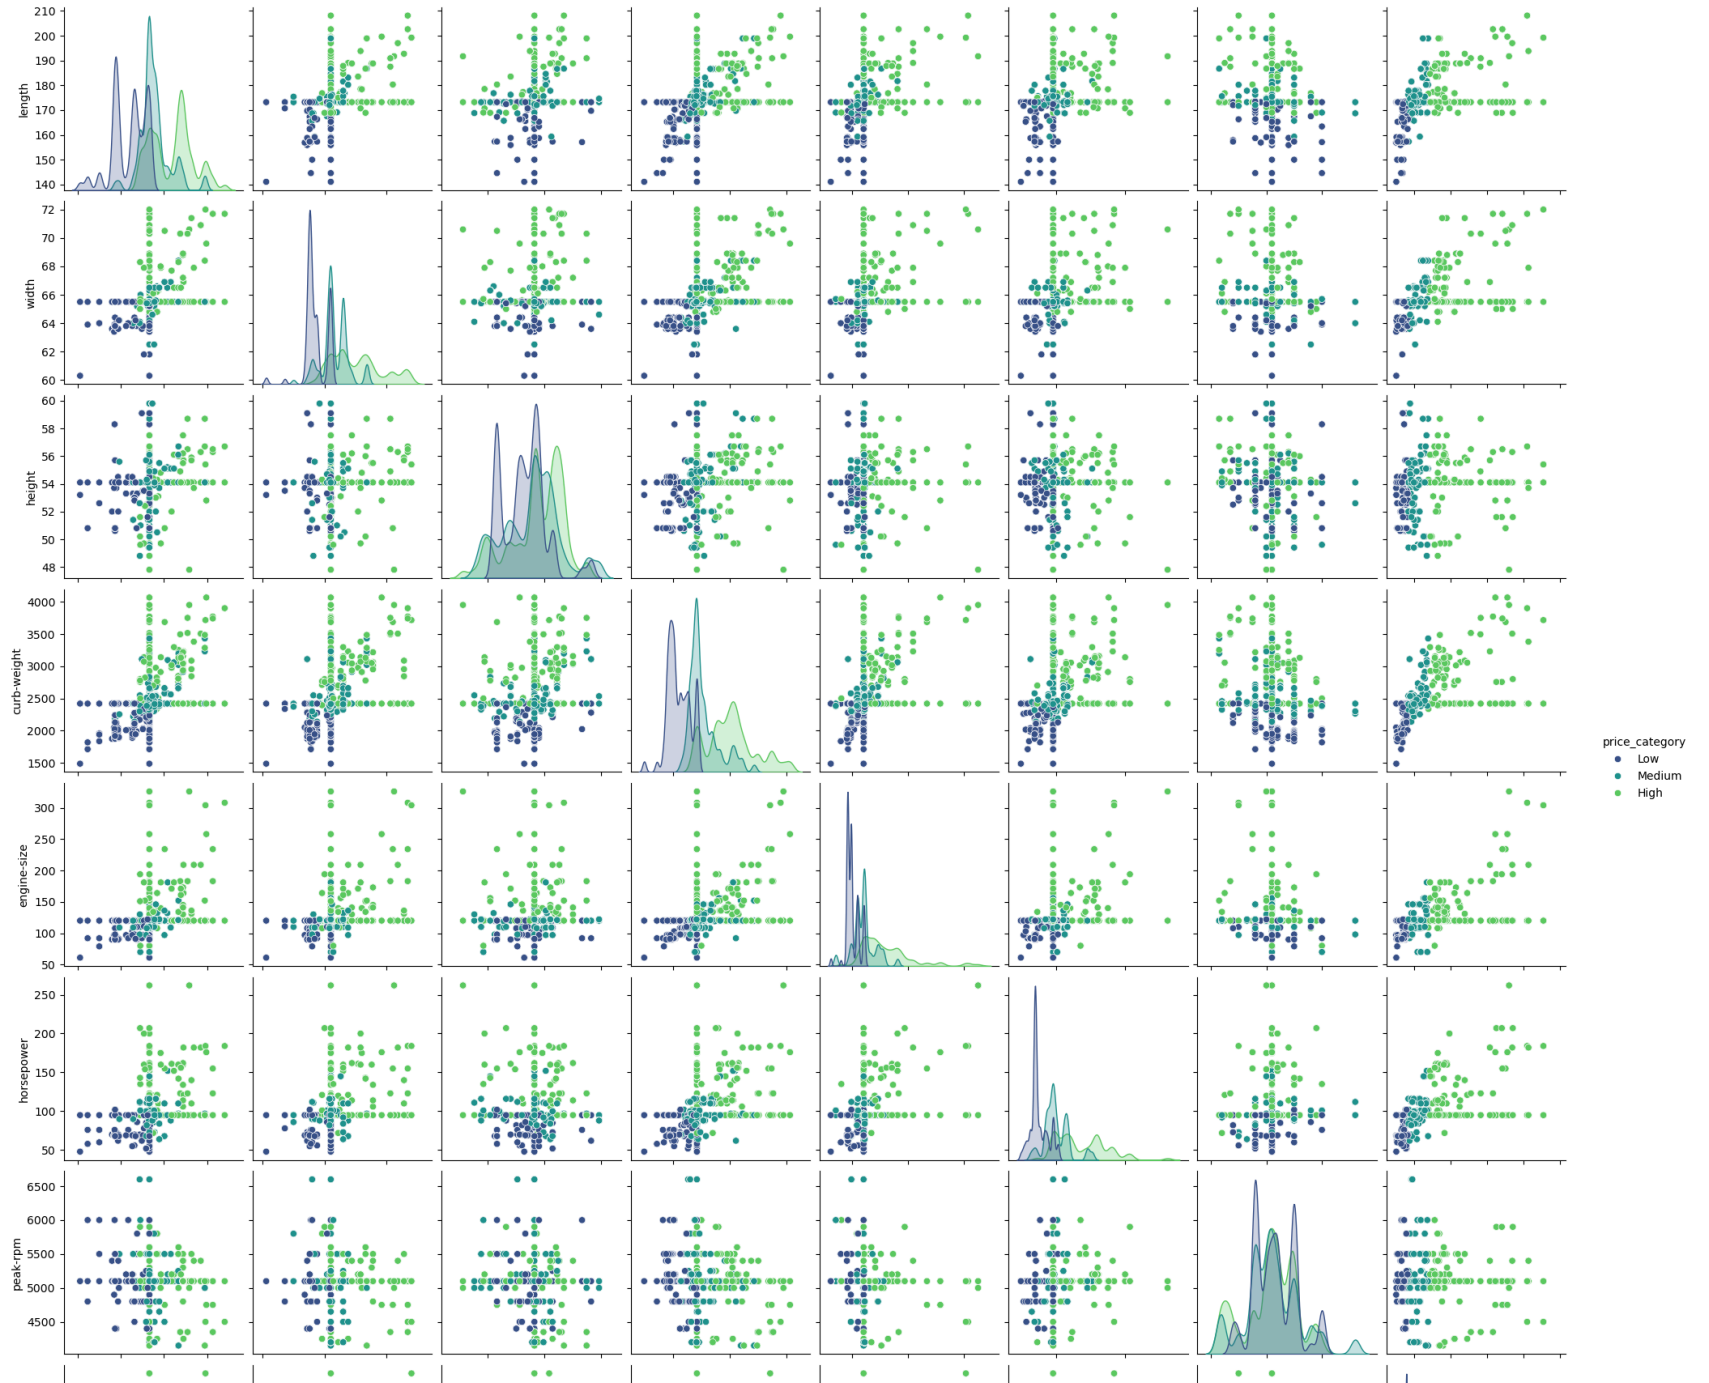
- Mercedes-Benz and Volvo are equipped with OHC and OHCV engine configurations. Among these, Mercedes-Benz exhibits the highest average price, indicating a premium positioning across both engine types.
- Jaguar and Alfa Romeo offer a diverse range of engine architectures, including OHC, OHCV, and DOHC. Notably, Jaguar demonstrates the highest pricing across these configurations, reflecting its focus on high-performance and luxury segments.
- Porsche models are available with OHC and OHCF engine types. Within the OHCF category, Porsche leads with the most expensive offerings, underscoring its engineering exclusivity and brand prestige.
- BMW, Mercury, and Audi all utilize the OHC engine type. BMW stands out with the highest pricing among this group, consistent with its reputation for delivering performance-oriented vehicles with advanced engineering.

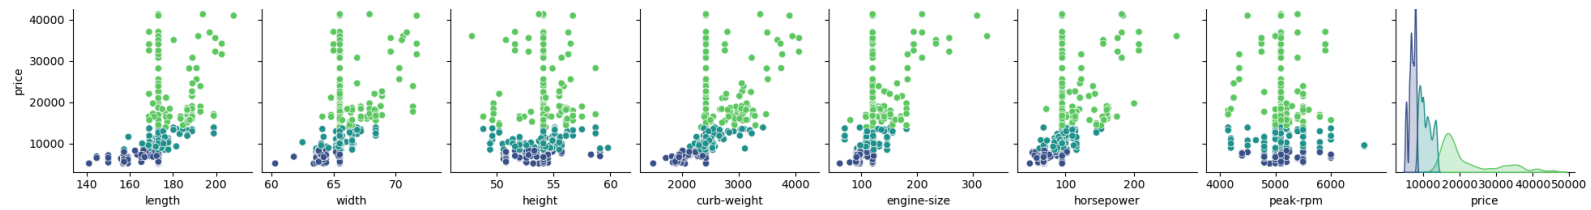
- Peugeot employs both OHC and inline engine types . Across both configurations, Peugeot leads in price within its engine categories, suggesting a relatively higher valuation for its engineering choices in the compact and mid-size car segments.

```
In [34]: df_subset = df[['length', 'width', 'height', 'curb-weight', 'engine-size', 'horsepower', 'peak-rpm', 'price']]
df_subset['price_category'] = pd.qcut(df_subset['price'], q=3, labels=['Low', 'Medium', 'High'])

sns.pairplot(df_subset, hue='price_category', palette='viridis')
plt.suptitle('Pairplot: Features vs. Price Category', y=1.02)
plt.show()
```

Pairplot: Features vs. Price Category

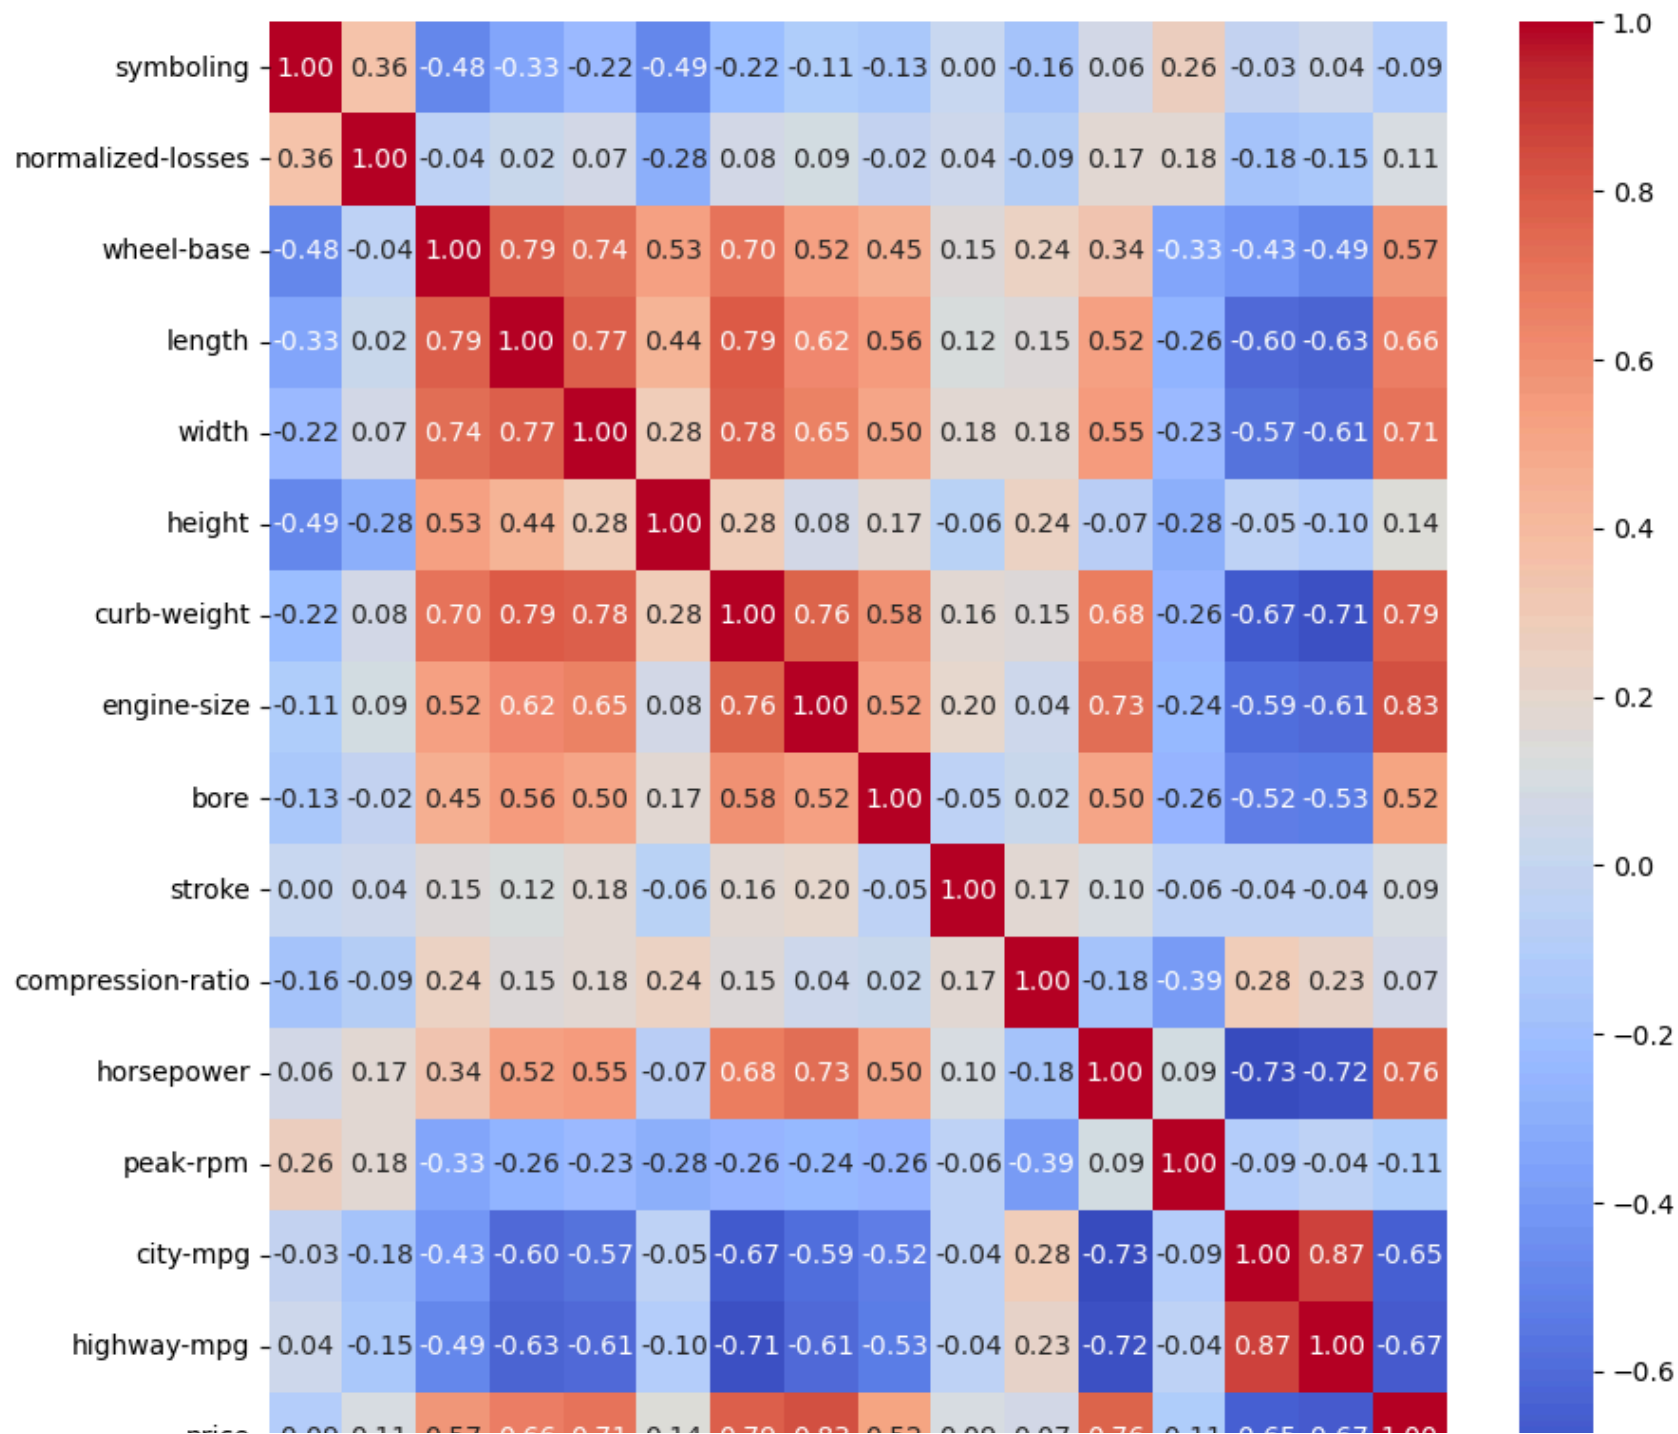


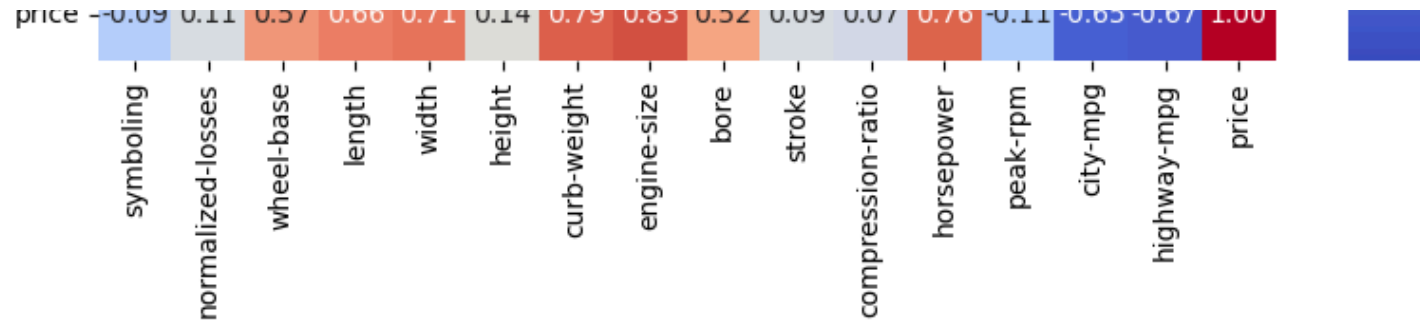


Correlation Matrix

```
In [35]: num = df.select_dtypes(include=['int64', 'float64']).columns

plt.figure(figsize=(10,10))
sns.heatmap(df[num].corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.show()
```





```
In [36]: correlation_with_target = df[numerical_features].corr()['price'].sort_values(ascending=False)

print("Top positively correlated:\n", correlation_with_target)
print("\nTop negatively correlated:\n", correlation_with_target)
```

Top positively correlated:

| | |
|-------------------|-----------|
| price | 1.000000 |
| engine-size | 0.827361 |
| curb-weight | 0.789834 |
| horsepower | 0.763740 |
| width | 0.709816 |
| length | 0.663409 |
| wheel-base | 0.566932 |
| bore | 0.517864 |
| height | 0.139918 |
| normalized-losses | 0.109223 |
| stroke | 0.092514 |
| compression-ratio | 0.073594 |
| symboling | -0.092012 |
| peak-rpm | -0.108704 |
| city-mpg | -0.650491 |
| highway-mpg | -0.665871 |

Name: price, dtype: float64

Top negatively correlated:

| | |
|-------------------|-----------|
| price | 1.000000 |
| engine-size | 0.827361 |
| curb-weight | 0.789834 |
| horsepower | 0.763740 |
| width | 0.709816 |
| length | 0.663409 |
| wheel-base | 0.566932 |
| bore | 0.517864 |
| height | 0.139918 |
| normalized-losses | 0.109223 |
| stroke | 0.092514 |
| compression-ratio | 0.073594 |
| symboling | -0.092012 |
| peak-rpm | -0.108704 |
| city-mpg | -0.650491 |
| highway-mpg | -0.665871 |

Name: price, dtype: float64

- **Target column: Price , engine-size, curb-weight, horsepower, width, length, wheel-base, bore, highway-mpg, city-mpg these are the highly correlated feature**

- **Engine size, curb-weight, horsepower, width highly correlated with target column so we keep them**

```
In [37]: features_df = df[numerical_features].drop(columns=['price'])
corr_matrix = features_df.corr()

threshold = 0.7

high_corr = (corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool)).stack().reset_index())
high_corr.columns = ['Feature 1', 'Feature 2', 'Correlation']
high_corr = high_corr[abs(high_corr['Correlation']) > threshold]

print("Highly correlated input features:\n", high_corr.sort_values(by='Correlation', ascending=False))
```

Highly correlated input features:

| | Feature 1 | Feature 2 | Correlation |
|-----|-------------|-------------|-------------|
| 104 | city-mpg | highway-mpg | 0.866629 |
| 41 | length | curb-weight | 0.792465 |
| 27 | wheel-base | length | 0.788264 |
| 51 | width | curb-weight | 0.775055 |
| 39 | length | width | 0.773800 |
| 69 | curb-weight | engine-size | 0.764232 |
| 28 | wheel-base | width | 0.735519 |
| 80 | engine-size | horsepower | 0.733774 |
| 30 | wheel-base | curb-weight | 0.704463 |
| 76 | curb-weight | highway-mpg | -0.711205 |
| 101 | horsepower | highway-mpg | -0.718016 |
| 100 | horsepower | city-mpg | -0.734102 |

- **Engine size, curb-weight, horsepower, highly correlated with target column**
- **Remove length, wheel-base, width, city-mpg**

Linear Regression

```
In [38]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.linear_model import LinearRegression
```



```
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
```

I have to apply the following ML techniques:

- To evaluate different machine learning algorithms on the dataset, I implemented and applied pipelines for three regression models: Linear Regression, Decision Tree Regressor, and Random Forest Regressor. Each model was encapsulated in a separate Pipeline, which included both preprocessing steps and the respective regressor. These pipelines were then fitted and evaluated using the same training and testing sets to ensure a fair comparison of model performance.

Linear Regression Pipeline

- I make a pipeline for linear regression
- Encode categorical variables, for the model
- One hot encode for nominal cat features, and Ordinal encode for ordinal cat features

```
In [39]: cylinders_order = ['two', 'three', 'four', 'five', 'six', 'eight', 'twelve']
doors_order = ['two', 'four']

lr_trans_1 = ColumnTransformer([

    ('ohe', OneHotEncoder(sparse_output=False, drop='first', handle_unknown='ignore'),
    ['make', 'engine-location', 'fuel-type', 'aspiration', 'body-style', 'drive-wheels', 'engine-type', 'fuel-system']),
    ('ord_enco', OrdinalEncoder(handle_unknown='use_encoded_value', unknown_value=-1, categories=[cylinders_order, doors_order]),
    ['num-of-cylinders', 'num-of-doors'])
], remainder='passthrough')
```

Find the less correlated features with the target variable and remove them

```
In [40]: features_df = df.drop(columns=['price'])

X_transformed = lr_trans_1.fit_transform(features_df)
X_transformed_df = pd.DataFrame(X_transformed, columns=lr_trans_1.get_feature_names_out())

corr_matrix = X_transformed_df.corr()
threshold = 0.7
```

```

high_corr = (corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool)).stack().reset_index())
high_corr.columns = ['Feature 1', 'Feature 2', 'Correlation']
high_corr = high_corr[abs(high_corr['Correlation']) > threshold]

print("Highly correlated input features:\n", high_corr.sort_values(by='Correlation', ascending=False))

```

Highly correlated input features:

| | Feature 1 | Feature 2 | Correlation |
|------|---------------------------|-----------------------------|-------------|
| 1496 | ohe_fuel-system_idi | remainder_compression-ratio | 0.879595 |
| 1710 | remainder_city-mpg | remainder_highway-mpg | 0.866629 |
| 647 | ohe_make_peugeot | ohe_engine-type_l | 0.851185 |
| 864 | ohe_make_subaru | ohe_engine-type_ohcf | 0.801845 |
| 1647 | remainder_length | remainder_curb-weight | 0.792465 |
| 1633 | remainder_wheel-base | remainder_length | 0.788264 |
| 1657 | remainder_width | remainder_curb-weight | 0.775055 |
| 1412 | ohe_engine-type_rotor | ohe_fuel-system_4bbl | 0.774690 |
| 1645 | remainder_length | remainder_width | 0.773800 |
| 1675 | remainder_curb-weight | remainder_engine-size | 0.764232 |
| 1583 | ord_enco_num-of-cylinders | remainder_engine-size | 0.746622 |
| 727 | ohe_make_porsche | ohe_engine-location_rear | 0.745585 |
| 1634 | remainder_wheel-base | remainder_width | 0.735519 |
| 1686 | remainder_engine-size | remainder_horsepower | 0.733774 |
| 315 | ohe_make_isuzu | ohe_fuel-system_spfi | 0.718352 |
| 1636 | remainder_wheel-base | remainder_curb-weight | 0.704463 |
| 1682 | remainder_curb-weight | remainder_highway-mpg | -0.711205 |
| 1707 | remainder_horsepower | remainder_highway-mpg | -0.718016 |
| 1706 | remainder_horsepower | remainder_city-mpg | -0.734102 |
| 1076 | ohe_fuel-type_gas | remainder_compression-ratio | -0.885288 |
| 1059 | ohe_fuel-type_gas | ohe_fuel-system_idi | -0.888635 |
| 1246 | ohe_drive-wheels_fwd | ohe_drive-wheels_rwd | -0.922378 |

```

In [41]: X1 = df.drop(['symboling', 'normalized-losses', 'bore', 'stroke', 'height', 'compression-ratio', 'symboling', 'peak-rpr',
y1 = df['price']]

```

```

In [42]: lr_pipeline_1 = Pipeline([('preprocessor', lr_trans_1),
('regressor', LinearRegression())])

```

```

X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, test_size=0.2, random_state=42)
lr_pipeline_1.fit(X_train1, y_train1)

```

```
y_train_pred1 = lr_pipeline_1.predict(X_train1)
y_test_pred1 = lr_pipeline_1.predict(X_test1)
```

Evaluation Metrics

```
In [43]: # MSE
train_mse1 = mean_squared_error(y_train1, y_train_pred1)
test_mse1 = mean_squared_error(y_test1, y_test_pred1)
# MAE
train_mae1 = mean_absolute_error(y_train1, y_train_pred1)
test_mae1 = mean_absolute_error(y_test1, y_test_pred1)
# RMSE
train_rmse1 = np.sqrt(mean_squared_error(y_train1, y_train_pred1))
test_rmse1 = np.sqrt(mean_squared_error(y_test1, y_test_pred1))
# R2 Score
train_r21 = r2_score(y_train1, y_train_pred1)
test_r21 = r2_score(y_test1, y_test_pred1)
# Adjusted R2
def adjusted_r21(r2, n, k):
    return 1 - (1 - r2) * ((n - 1) / (n - k - 1))
n_train1, k1 = X_train1.shape
n_test1 = X_test1.shape[0]
train_adj_r21 = adjusted_r21(train_r21, n_train1, k1)
test_adj_r21 = adjusted_r21(test_r21, n_test1, k1)

print('Linear Regression Model')
print(f"\nTrain MSE: {train_mse1:.4f}")
print(f"Test MSE: {test_mse1:.4f}")
print('=====')
print(f"\nTrain MAE: {train_mae1:.4f}")
print(f"Test MAE: {test_mae1:.4f}")
print('=====')
print(f"\nTrain RMSE: {train_rmse1:.4f}")
print(f"Test RMSE: {test_rmse1:.4f}")
print('=====')
print(f"\nTrain R2: {train_r21:.4f}")
print(f"Test R2: {test_r21:.4f}")
print('=====')
print(f"\nTrain Adjusted R2: {train_adj_r21:.4f}")
print(f"Test Adjusted R2: {test_adj_r21:.4f}")
```

Linear Regression Model

Train MSE: 5162039.7729

Test MSE: 5254884.5825

=====

Train MAE: 1537.4394

Test MAE: 1565.8101

=====

Train RMSE: 2272.0123

Test RMSE: 2292.3535

=====

Train R²: 0.9188

Test R²: 0.9154

=====

Train Adjusted R²: 0.9187

Test Adjusted R²: 0.9152

- Compared to the model with all features, we achieved similar performance using fewer features, while still maintaining a good fit without signs of overfitting

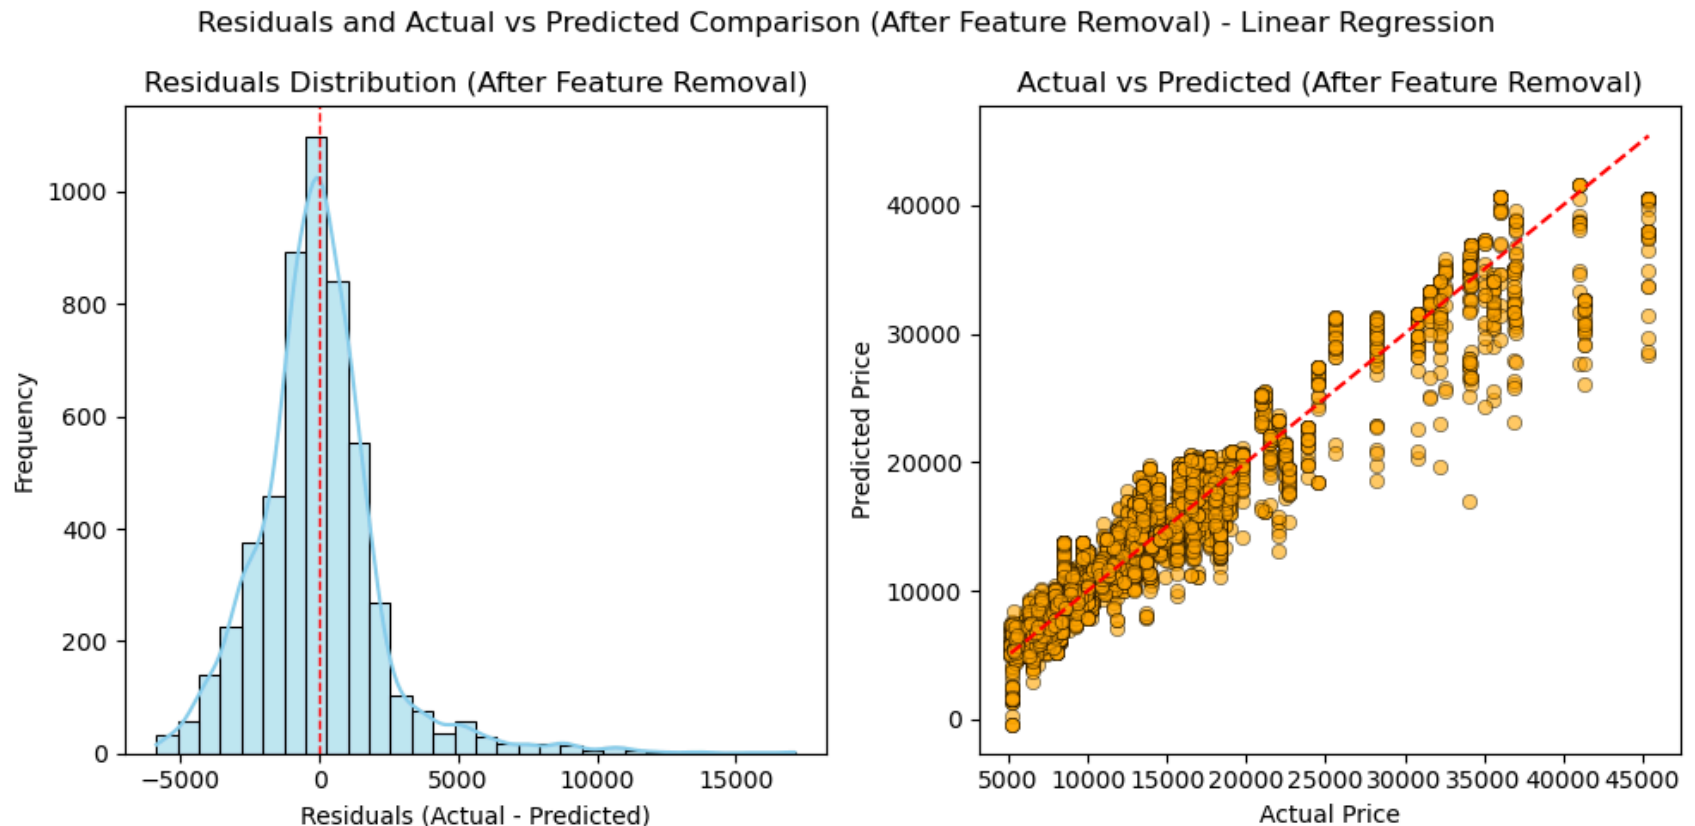
Visualization of residuals to evaluate model performance for both feature sets

```
In [44]: fig, ax = plt.subplots(1, 2, figsize=(10, 5))

residuals1 = y_test1 - y_test_pred1
sns.histplot(residuals1, bins=30, kde=True, color="skyblue", edgecolor="black", ax=ax[0])
ax[0].axvline(x=0, color='red', linestyle='--', linewidth=1)
ax[0].set_title("Residuals Distribution (After Feature Removal)")
ax[0].set_xlabel("Residuals (Actual - Predicted)")
ax[0].set_ylabel("Frequency")

sns.scatterplot(x=y_test1, y=y_test_pred1, alpha=0.6, color='orange', edgecolor='k', ax=ax[1])
ax[1].plot([y_test1.min(), y_test1.max()], [y_test1.min(), y_test1.max()], color='red', linestyle='--')
ax[1].set_title("Actual vs Predicted (After Feature Removal)")
ax[1].set_xlabel("Actual Price")
ax[1].set_ylabel("Predicted Price")
```

```
plt.suptitle("Residuals and Actual vs Predicted Comparison (After Feature Removal) - Linear Regression")
plt.tight_layout()
plt.show()
```



Decision Tree Regression

```
In [45]: from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.model_selection import GridSearchCV
```

Pipeline for Decision Tree Regression

```
In [46]: cylinders_order = ['two', 'three', 'four', 'five', 'six', 'eight', 'twelve']
doors_order = ['two', 'four']

dtr1 = ColumnTransformer([
    ('ohe', OneHotEncoder(sparse_output=False, drop='first', handle_unknown='ignore'),
    ['make', 'engine-location', 'fuel-type', 'aspiration', 'body-style', 'drive-wheels', 'engine-type', 'fuel-system'],
    ('ordenco', OrdinalEncoder(handle_unknown='use_encoded_value', unknown_value=-1,
                                categories=[cylinders_order, doors_order])),
    ['num-of-cylinders', 'num-of-doors'])
], remainder='passthrough')
```

```
In [47]: dtrx_1 = df.drop(['symboling', 'normalized-losses', 'bore', 'stroke', 'height',
                        'compression-ratio', 'peak-rpm', 'highway-mpg', 'price'], axis=1)
dtry_1 = df['price']
```

Best Hyperparameters Found by GridSearchCV

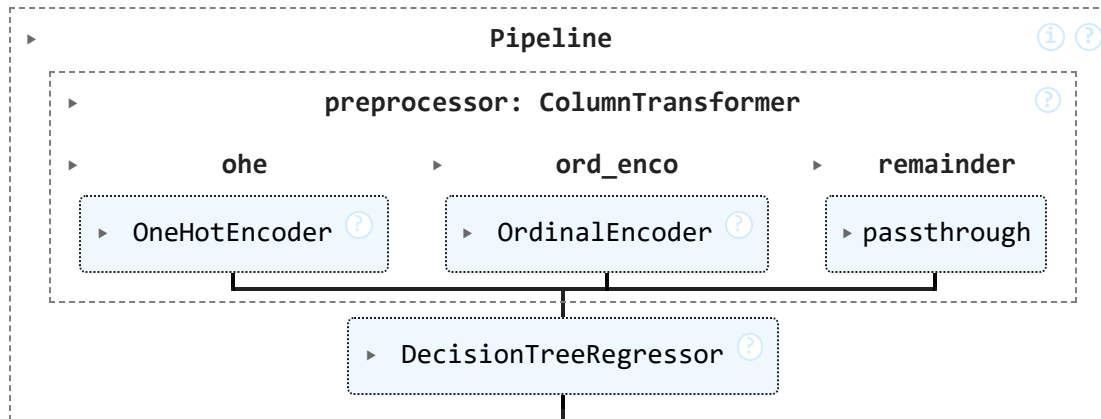
- {'regressor__ccp_alpha': 0.0001, 'regressor__max_depth': 40, 'regressor__max_features': 25, 'regressor__max_leaf_nodes': 100, 'regressor__min_samples_leaf': 2, 'regressor__min_samples_split': 10, 'regressor__splitter': 'random'}

```
In [48]: dtr1_pipeline = Pipeline([
    ('preprocessor', dtr1),
    ('regressor', DecisionTreeRegressor(ccp_alpha=0.0001, max_depth=40, max_features=25,
                                         max_leaf_nodes=100, min_samples_leaf=2, min_samples_split=10, splitter='random'))

X_train_dtr, X_test_dtr, y_train_dtr, y_test_dtr = train_test_split(dtrx_1, dtry_1, test_size=0.2, random_state=42)

dtr1_pipeline.fit(X_train_dtr, y_train_dtr)
```

Out[48]:



Evaluation Metrics for DTR

```
In [49]: y_train_pred_dtr = dtr1_pipeline.predict(X_train_dtr)
y_test_pred_dtr = dtr1_pipeline.predict(X_test_dtr)

# MSE
train_mse_dtr = mean_squared_error(y_train_dtr, y_train_pred_dtr)
test_mse_dtr = mean_squared_error(y_test_dtr, y_test_pred_dtr)
# MAE
train_mae_dtr = mean_absolute_error(y_train_dtr, y_train_pred_dtr)
test_mae_dtr = mean_absolute_error(y_test_dtr, y_test_pred_dtr)
# RMSE
train_rmse_dtr = np.sqrt(mean_squared_error(y_train_dtr, y_train_pred_dtr))
test_rmse_dtr = np.sqrt(mean_squared_error(y_test_dtr, y_test_pred_dtr))
# R²
train_r2_dtr = r2_score(y_train_dtr, y_train_pred_dtr)
test_r2_dtr = r2_score(y_test_dtr, y_test_pred_dtr)
# Adjusted R²
def adjusted_r2(r2, n, k):
    return 1 - (1 - r2) * ((n - 1) / (n - k - 1))
n_train_dtr, k_dtr = X_train_dtr.shape
n_test_dtr = X_test_dtr.shape[0]
train_adj_r2_dtr = adjusted_r2(train_r2_dtr, n_train_dtr, k_dtr)
test_adj_r2_dtr = adjusted_r2(test_r2_dtr, n_test_dtr, k_dtr)

print('Decision Tree Regressor')
print(f"\nTrain MSE: {train_mse_dtr:.4f}")
print(f"Test MSE: {test_mse_dtr:.4f}")
```

```

print('=====')
print(f"\nTrain MAE: {train_mae_dtr:.4f}")
print(f"Test MAE: {test_mae_dtr:.4f}")
print('=====')
print(f"\nTrain RMSE: {train_rmse_dtr:.4f}")
print(f"Test RMSE: {test_rmse_dtr:.4f}")
print('=====')
print(f"\nTrain R²: {train_r2_dtr:.4f}")
print(f"Test R²: {test_r2_dtr:.4f}")
print('=====')
print(f"\nTrain Adjusted R²: {train_adj_r2_dtr:.4f}")
print(f"Test Adjusted R²: {test_adj_r2_dtr:.4f}")

```

Decision Tree Regressor

Train MSE: 1715425.7840

Test MSE: 1776206.9677

=====

Train MAE: 853.7516

Test MAE: 857.2577

=====

Train RMSE: 1309.7426

Test RMSE: 1332.7441

=====

Train R²: 0.9730

Test R²: 0.9714

=====

Train Adjusted R²: 0.9730

Test Adjusted R²: 0.9713

- After removing less important features, the Decision Tree Regressor showed balanced performance with minimal overfitting.
- Both Train and Test R² remained high (0.97), and error metrics like MAE and RMSE were nearly identical, indicating strong generalization

```

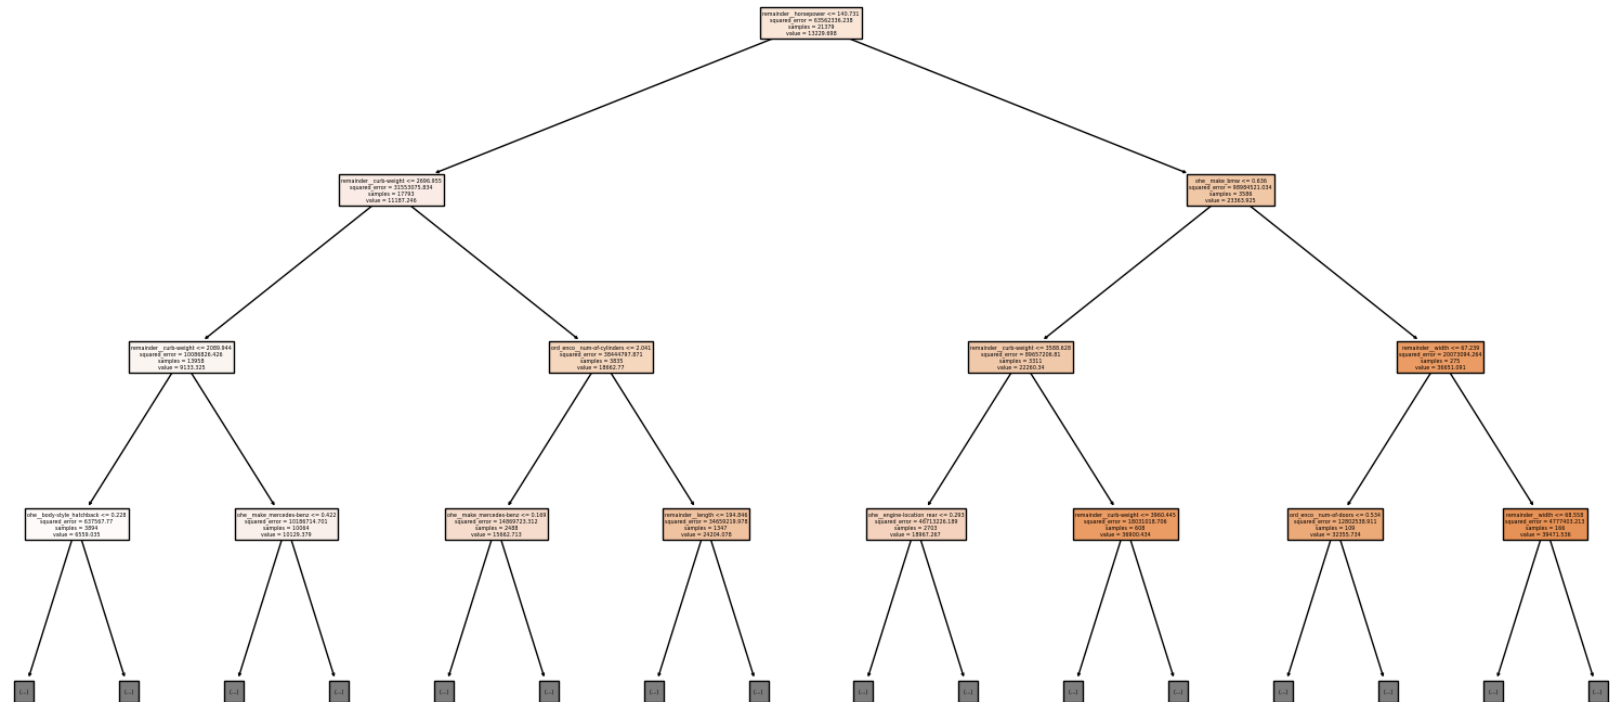
In [50]: preprocessor = dtr1_pipeline.named_steps['preprocessor']
feature_names = preprocessor.get_feature_names_out()

```



```
In [51]: tree_model = dtr1_pipeline.named_steps['regressor']
feature_names = dtr1_pipeline.named_steps['preprocessor'].get_feature_names_out()

plt.figure(figsize=(20, 10))
plot_tree(tree_model, filled=True, feature_names=feature_names, max_depth=3)
plt.show()
```



```
In [52]: fig, ax = plt.subplots(1, 2, figsize=(10, 5))

residualsdtr = y_test_dtr - y_test_pred_dtr
sns.histplot(residualsdtr, bins=30, kde=True, color="skyblue", edgecolor="black", ax=ax[0])
ax[0].axvline(x=0, color='red', linestyle='--', linewidth=1)
ax[0].set_title("Residuals Distribution (After Feature Removal)")
ax[0].set_xlabel("Residuals (Actual - Predicted)")
ax[0].set_ylabel("Frequency")
```

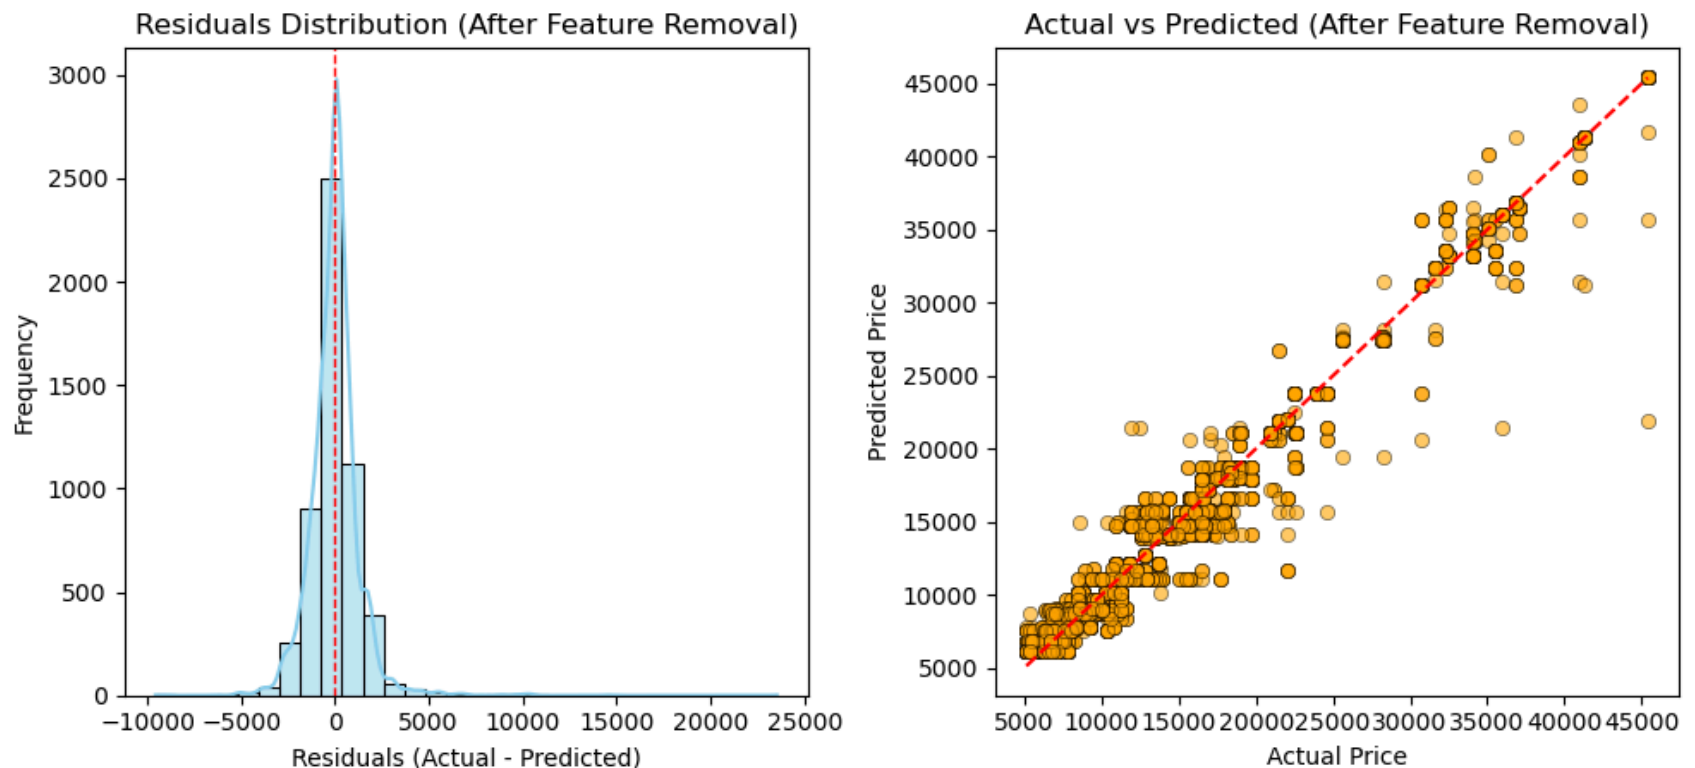
```

sns.scatterplot(x=y_test_dtr, y=y_test_pred_dtr, alpha=0.6, color='orange', edgecolor='k', ax=ax[1])
ax[1].plot([y_test_dtr.min(), y_test_dtr.max()], [y_test_dtr.min(), y_test_dtr.max()], color='red', linestyle='--')
ax[1].set_title("Actual vs Predicted (After Feature Removal)")
ax[1].set_xlabel("Actual Price")
ax[1].set_ylabel("Predicted Price")

plt.suptitle("Residuals and Actual vs Predicted Comparison (After Feature Removal) - Decision Tree Regressor")
plt.tight_layout()
plt.show()

```

Residuals and Actual vs Predicted Comparison (After Feature Removal) - Decision Tree Regressor



Random Forest Regressor

```
In [ ]: from sklearn.ensemble import RandomForestRegressor
```

```
In [54]: cylinders_order = ['two', 'three', 'four', 'five', 'six', 'eight', 'twelve']
doors_order = ['two', 'four']

rf_transformer1 = ColumnTransformer([
    ('ohe', OneHotEncoder(sparse_output=False, drop='first', handle_unknown='ignore'),
    ['make', 'engine-location', 'fuel-type', 'aspiration', 'body-style', 'drive-wheels', 'engine-type', 'fuel-system'],
    ('ord', OrdinalEncoder(handle_unknown='use_encoded_value', unknown_value=-1,
                           categories=[cylinders_order, doors_order])),
    ['num-of-cylinders', 'num-of-doors'])
], remainder='passthrough')
```

Best Parameters for Random Forest Regressor

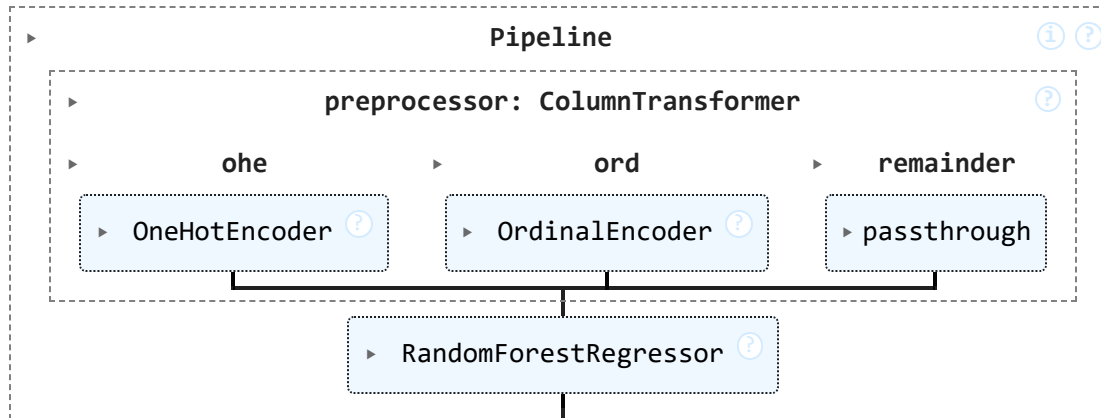
- Best Parameters: {'regressor__bootstrap': True, 'regressor__max_depth': 40, 'regressor__min_samples_leaf': 2, 'regressor__min_samples_split': 2, 'regressor__n_estimators': 100}

```
In [55]: rf_X1 = df.drop(['symboling', 'normalized-losses', 'bore', 'stroke', 'height',
                        'compression-ratio', 'peak-rpm', 'highway-mpg', 'price'], axis=1)
rf_y1 = df['price']

rf_pipeline1 = Pipeline([
    ('preprocessor', rf_transformer1),
    ('regressor', RandomForestRegressor(bootstrap=True, max_depth=40, min_samples_leaf=2,
                                       min_samples_split=2, n_estimators=100, random_state=42))])

X_train_rf1, X_test_rf1, y_train_rf1, y_test_rf1 = train_test_split(rf_X1, rf_y1, test_size=0.2, random_state=42)
rf_pipeline1.fit(X_train_rf1, y_train_rf1)
```

Out[55]:



```
In [56]: y_train_pred_rf1 = rf_pipeline1.predict(X_train_rf1)
y_test_pred_rf1 = rf_pipeline1.predict(X_test_rf1)

# MSE
train_mse_rf1 = mean_squared_error(y_train_rf1, y_train_pred_rf1)
test_mse_rf1 = mean_squared_error(y_test_rf1, y_test_pred_rf1)
# MAE
train_mae_rf1 = mean_absolute_error(y_train_rf1, y_train_pred_rf1)
test_mae_rf1 = mean_absolute_error(y_test_rf1, y_test_pred_rf1)
# RMSE
train_rmse_rf1 = np.sqrt(mean_squared_error(y_train_rf1, y_train_pred_rf1))
test_rmse_rf1 = np.sqrt(mean_squared_error(y_test_rf1, y_test_pred_rf1))
# R²
train_r2_rf1 = r2_score(y_train_rf1, y_train_pred_rf1)
test_r2_rf1 = r2_score(y_test_rf1, y_test_pred_rf1)
# Adjusted R² function
def adjusted_r2(r2, n, k):
    return 1 - (1 - r2) * ((n - 1) / (n - k - 1))
n_train_rf1, k_rf1 = X_train_rf1.shape
n_test_rf1 = X_test_rf1.shape[0]
train_adj_r2_rf1 = adjusted_r2(train_r2_rf1, n_train_rf1, k_rf1)
test_adj_r2_rf1 = adjusted_r2(test_r2_rf1, n_test_rf1, k_rf1)

print("Random Forest Model")
print(f"\nTrain MSE: {train_mse_rf1:.4f}")
print(f"Test MSE: {test_mse_rf1:.4f}")
print('=====')
print(f"\nTrain MAE: {train_mae_rf1:.4f}")
```

```

print(f"Test MAE: {test_mae_rf1:.4f}")
print('=====')
print(f"\nTrain RMSE: {train_rmse_rf1:.4f}")
print(f"Test RMSE: {test_rmse_rf1:.4f}")
print('=====')
print(f"\nTrain R²: {train_r2_rf1:.4f}")
print(f"Test R²: {test_r2_rf1:.4f}")
print('=====')
print(f"\nTrain Adjusted R²: {train_adj_r2_rf1:.4f}")
print(f"Test Adjusted R²: {test_adj_r2_rf1:.4f}")

```

Random Forest Model

Train MSE: 146727.3371

Test MSE: 315051.7191

=====

Train MAE: 113.1672

Test MAE: 153.1049

=====

Train RMSE: 383.0500

Test RMSE: 561.2947

=====

Train R²: 0.9977

Test R²: 0.9949

=====

Train Adjusted R²: 0.9977

Test Adjusted R²: 0.9949

```

In [57]: fig, ax = plt.subplots(1, 2, figsize=(10, 5))

residualsrfr = y_test_rf1 - y_test_pred_rf1
sns.histplot(residualsrfr, bins=30, kde=True, color="skyblue", edgecolor="black", ax=ax[0])
ax[0].axvline(x=0, color='red', linestyle='--', linewidth=1)
ax[0].set_title("Residuals Distribution (After Feature Removal)")
ax[0].set_xlabel("Residuals (Actual - Predicted)")
ax[0].set_ylabel("Frequency")

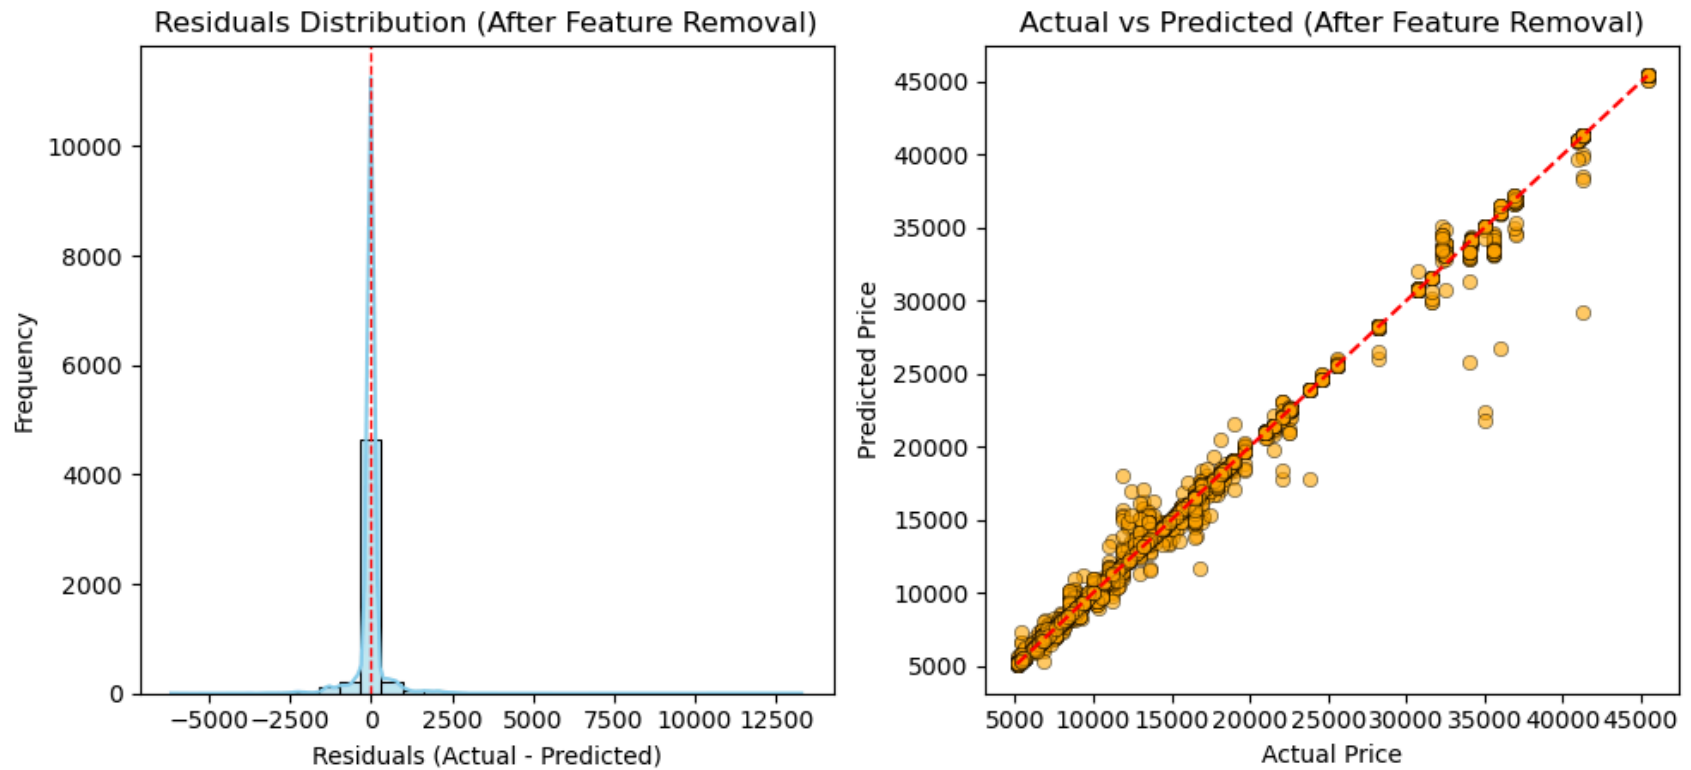
sns.scatterplot(x=y_test_rf1, y=y_test_pred_rf1, alpha=0.6, color='orange', edgecolor='k', ax=ax[1])

```

```
ax[1].plot([y_test_rf1.min(), y_test_rf1.max()], [y_test_rf1.min(), y_test_rf1.max()], color='red', linestyle='--')
ax[1].set_title("Actual vs Predicted (After Feature Removal)")
ax[1].set_xlabel("Actual Price")
ax[1].set_ylabel("Predicted Price")

plt.suptitle("Residuals and Actual vs Predicted Comparison (After Feature Removal) - Random Forest Regressor")
plt.tight_layout()
plt.show()
```

Residuals and Actual vs Predicted Comparison (After Feature Removal) - Random Forest Regressor



Comparison of MSE, R2 score for all three models

```
In [58]: metrics_lr = {'MSE': [train_mse1, test_mse1], 'R²': [train_r21, test_r21]}
metrics_dtr = {'MSE': [train_mse_dtr, test_mse_dtr], 'R²': [train_r2_dtr, test_r2_dtr]}
metrics_rf = {'MSE': [train_mse_rf1, test_mse_rf1], 'R²': [train_r2_rf1, test_r2_rf1]}

models = ['Linear Regression', 'Decision Tree Regressor', 'Random Forest']
```

```

metrics_combined = {
    'Linear Regression': [train_mse1, test_mse1, train_r21, test_r21],
    'Decision Tree Regressor': [train_mse_dtr, test_mse_dtr, train_r2_dtr, test_r2_dtr],
    'Random Forest': [train_mse_rf1, test_mse_rf1, train_r2_rf1, test_r2_rf1]}

fig, ax = plt.subplots(1, 2, figsize=(14, 6))

bar_width = 0.3
x_pos = np.arange(len(models))

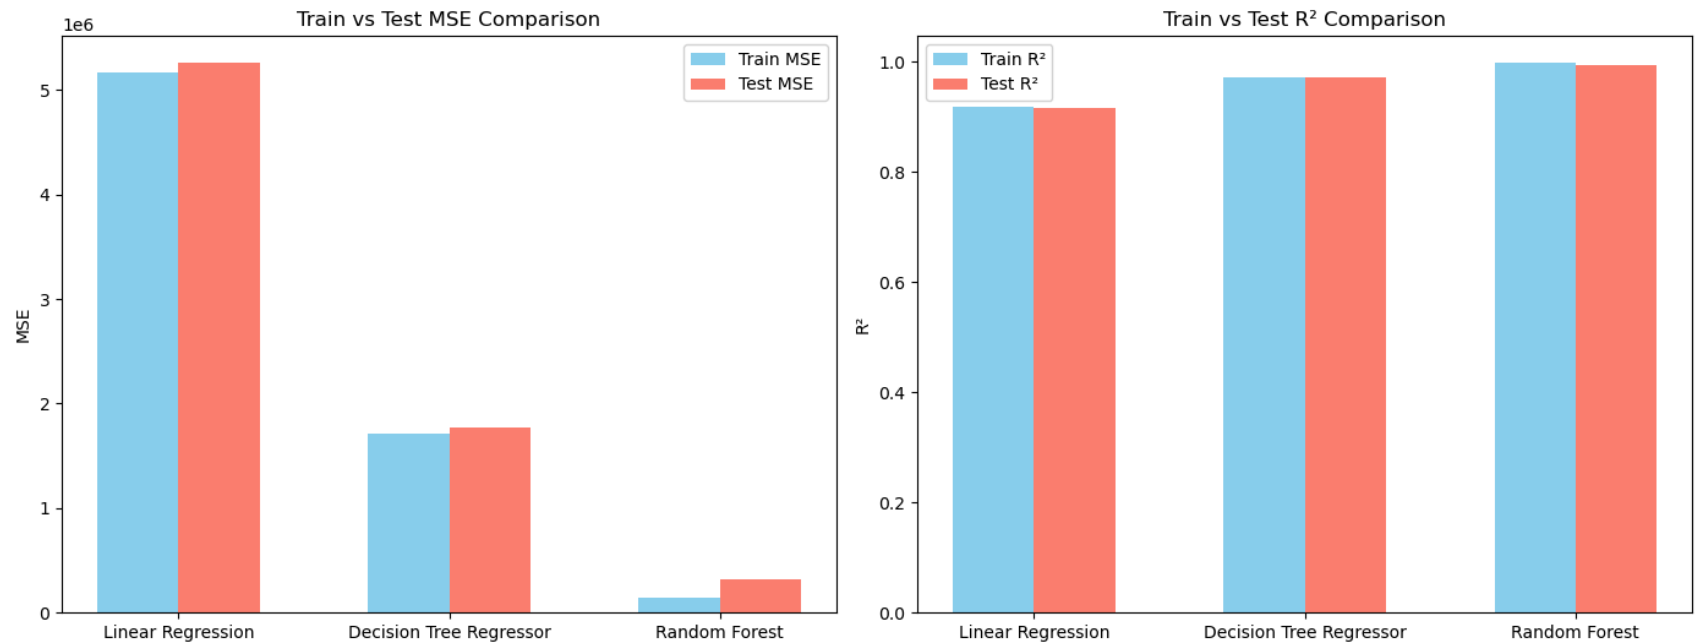
ax[0].bar(x_pos - bar_width/2, [metrics_combined[model][0] for model in models], bar_width, label='Train MSE', color='red')
ax[0].bar(x_pos + bar_width/2, [metrics_combined[model][1] for model in models], bar_width, label='Test MSE', color='blue')
ax[0].set_title('Train vs Test MSE Comparison')
ax[0].set_ylabel('MSE')
ax[0].set_xticks(x_pos)
ax[0].set_xticklabels(models)
ax[0].legend()

ax[1].bar(x_pos - bar_width/2, [metrics_combined[model][2] for model in models], bar_width, label='Train R2', color='red')
ax[1].bar(x_pos + bar_width/2, [metrics_combined[model][3] for model in models], bar_width, label='Test R2', color='blue')
ax[1].set_title('Train vs Test R2 Comparison')
ax[1].set_ylabel('R2')
ax[1].set_xticks(x_pos)
ax[1].set_xticklabels(models)
ax[1].legend()

plt.suptitle('Comparison of MSE and R2 for Models (Train vs Test)', fontsize=16)
plt.tight_layout()
plt.subplots_adjust(top=0.85)
plt.show()

```

Comparison of MSE and R² for Models (Train vs Test)



Model Comparison Summary

- The Random Forest model stands out as the best performer, with the highest R² and adjusted R² values, indicating excellent predictive power. It also has the lowest MAE and RMSE on the training set, making it highly accurate.
- Decision Tree Regressor also performs exceptionally well, with high R² values and relatively low MSE and RMSE compared to Linear Regression.
- Linear Regression, while a strong model with decent results, falls short in comparison to the more complex models, showing higher MSE, MAE, and RMSE, and a slightly lower R².

Recommendation

- If computational efficiency and simplicity are key, Linear Regression could be used, but for better accuracy and predictive power, the Decision Tree Regressor or Random Forest models should be prioritized. Random Forest, in particular, should be considered the best option for this dataset based on its strong performance across all metrics.

Reporting & Insights

What key insights did you gain from EDA about car prices

- Car Brand Popularity: Toyota is the most popular car brand, making up about 24% of the dataset. This suggests Toyota has a strong market presence or is more widely available in the dataset.
- Car Type Distribution: Sedans represent the majority of cars in the dataset, accounting for 50%. This aligns with consumer preferences for vehicles offering comfort, fuel efficiency, and practicality.
- Engine Type Preference: The most common engine types in the dataset are the OHC (Overhead Camshaft) engine, found in 70% of the cars, and the 4-cylinder engine, which makes up 78% of the entries. These engine types offer a balance between performance and fuel efficiency, which could explain their popularity.
- Fuel System Trends: Around 50% of the cars use the MPFI (Multi-Point Fuel Injection) system, reflecting a modern trend towards improving fuel efficiency and reducing emissions.
- Drive Type: Front-wheel drive (FWD) is the most common configuration, present in 60% of the cars. This could reflect the cost-effectiveness and fuel efficiency of FWD vehicles.
- Engine Location: 95% of the cars have a front engine location, which is typical for most car models and configurations.
- Price Variations by Features:
 - Luxury and high-performance cars, such as those from Jaguar, Mercedes-Benz, Porsche, and BMW, tend to have significantly higher prices.
 - Diesel cars and those with turbocharged engines also tend to have higher prices, likely due to the more advanced engine technology and higher performance.
 - RWD (Rear-Wheel Drive) cars tend to be more expensive, as they are often associated with premium, sports, and luxury vehicles.

Which features had the most impact on price prediction

- Car Brand
- Engine Type
- Fuel System
- Drive Type
- Engine Configuration
- Body Style
- Horsepower
- MPG

What challenges did you face during preprocessing and modeling

- The data had some anomalies like (wrong data type, wrong values, incorrect data format) etc, missing values, and outliers.
- In modeling we had to deal with the categorical data, so we had to encode it, then find the highly correlated features with target variable, then remove the highly correlated input features to input features, then some feature selection.
- Too much time spent on model tuning.

If given a larger dataset with more features, what additional steps would you take?

- Feature Engineering: With more data, I would explore new features that might better capture the relationships.
- Model Optimization: With a larger dataset, I would experiment with more complex models like ensemble methods (Random Forest, Gradient Boosting).
- Address Multicollinearity: I would use techniques like Principal Component Analysis (PCA) or regularization (L1/L2 regularization) to handle multicollinearity.